

# ML Cheat Sheet

## 1 Math Prerequisites

### 1.1 Derivatives

- $\partial(\mathbf{XY}) = (\partial\mathbf{X})\mathbf{Y} + \mathbf{X}(\partial\mathbf{Y})$
- $\frac{\partial \mathbf{f}(\mathbf{g}(\mathbf{u}(\mathbf{x})))}{\partial \mathbf{x}} = \frac{\partial \mathbf{u}(\mathbf{x})}{\partial \mathbf{x}} \frac{\partial \mathbf{g}(\mathbf{u})}{\partial \mathbf{u}} \frac{\partial \mathbf{f}(\mathbf{g})}{\partial \mathbf{g}}$
- $\frac{\partial \mathbf{x}^T \mathbf{a}}{\partial \mathbf{x}} = \frac{\partial \mathbf{a}^T \mathbf{x}}{\partial \mathbf{x}} = \mathbf{a}$
- $\frac{\partial \mathbf{a}^T \mathbf{X} \mathbf{b}}{\partial \mathbf{X}} = \mathbf{ab}^T$
- $\frac{\partial \mathbf{a}^T \mathbf{X}^T \mathbf{b}}{\partial \mathbf{X}} = \mathbf{ba}^T$
- $\frac{\partial \mathbf{x}}{\partial \mathbf{X}_{ij}} = \mathbf{J}^{ij}, \mathbf{J}^{ij}$  is the single entry matrix
- $\frac{\partial \mathbf{b}^T \mathbf{X}^T \mathbf{X} \mathbf{c}}{\partial \mathbf{X}} = \mathbf{X}(\mathbf{bc}^T + \mathbf{cb}^T)$
- $\frac{\partial \mathbf{x}^T \mathbf{B} \mathbf{x}}{\partial \mathbf{x}} = (\mathbf{B} + \mathbf{B}^T) \mathbf{x}$
- $\frac{\partial}{\partial \mathbf{x}} (\mathbf{x} - \mathbf{As})^T \mathbf{W}(\mathbf{x} - \mathbf{As}) = 2\mathbf{W}(\mathbf{x} - \mathbf{As})$
- $\frac{\partial}{\partial \mathbf{X}} \|\mathbf{X}\|_F^2 = \frac{\partial}{\partial \mathbf{X}} \text{Tr}(\mathbf{XX}^H) = 2\mathbf{X}$

### 1.2 Linear Algebra

- **positive definite** (pd) if  $\mathbf{a}^T \mathbf{V} \mathbf{a} > 0$
- $(\mathbf{x} - \mathbf{b})^T (\mathbf{x} - \mathbf{b}) = \|\mathbf{x} - \mathbf{b}\|_2^2$
- $\|\mathbf{X}\|_F = \|\mathbf{X}^T\|_F$

### 1.3 Distributions

Valid distribution  $p(x) > 0, \forall x$  and  $\sum p(x) = 1$   
Model is identifiable iff  $\theta_1 = \theta_2 \rightarrow P_{\theta_1} = P_{\theta_2}$

- **Gaussian** (Not convex):

$$\mathcal{N}(x|\mu, \sigma^2) = \frac{1}{\sqrt{2\pi\sigma^2}} \exp\left(-\frac{(x-\mu)^2}{2\sigma^2}\right)$$

$$\mathcal{N}(x|\mu, \Sigma^2) = \frac{\exp\left(-\frac{1}{2}(\mathbf{x}-\mu)^T \Sigma^{-1}(\mathbf{x}-\mu)\right)}{\sqrt{(2\pi)^D \det(\Sigma)}}$$

- **Poisson**:  $P(k \text{ events in interval}) = e^{-\lambda} \frac{\lambda^k}{k!}$
- **Bernoulli**:  $p(y|\mu) = \mu^y (1-\mu)^{1-y}$

### 1.4 Convexity

A function  $f(x)$  is convex if

- for any  $\mathbf{x}_1, \mathbf{x}_2 \in \mathbf{X}$  and  $0 \leq \lambda \leq 1$ , we have :  
 $f(\lambda \mathbf{x}_1 + (1-\lambda)\mathbf{x}_2) \leq \lambda f(\mathbf{x}_1) + (1-\lambda)f(\mathbf{x}_2)$
- it is a sum of convex functions
- composition of convex and linear functions
- $f(x) = g(h(x))$ ,  $g, h$  are convex,  $g$  increasing
- the Hessian  $\mathbf{H}$  is positive semi-definite

### 1.5 Others

- Production of independent variables:  
 $\text{Var}(xy) = \mathbb{E}(x^2) \mathbb{E}(y^2) - [\mathbb{E}(x)]^2 [\mathbb{E}(y)]^2$
- Covariance matrix of a data vector  $\mathbf{x}$

$$\Sigma = \frac{1}{N} \sum_{n=1}^N (\mathbf{x}_n - \mathbb{E}(\mathbf{x}))(\mathbf{x}_n - \mathbb{E}(\mathbf{x}))^T$$

- Multi-class  $\mathbf{x}$

$$p(\mathbf{y}|\mathbf{X}, \beta) = \prod p(\mathbf{y}_n|\mathbf{x}_n, \beta)$$

$$= \prod \prod [p(\mathbf{y}_n = k|\mathbf{x}_n, \beta)]^{y_{nk}}$$

## 2 Cost functions

**Mean square error (MSE):**

$$MSE(\mathbf{w}) = \frac{1}{N} \sum_{n=1}^N (y_n - f(\mathbf{x}_n))^2$$

- MSE is **strictly convex** thus it has only one global minimum value.
- MSE is very prone to outliers.

**Mean Absolute Error (MAE):**

$$MAE = \frac{1}{N} \sum_{n=1}^N |y_n - f(\mathbf{x}_n)|$$

- MAE is more robust to outliers than MSE.

**Huber loss**

$$Huber = \begin{cases} \frac{1}{2} z^2 & , |z| \leq \delta \\ \delta |z| - \frac{1}{2} \delta^2 & , |z| > \delta \end{cases}$$

- Huber loss is convex, differentiable, and also robust to outliers but hard to set  $\delta$ .

**Tukey's bisquare loss**

$$L(z) = \begin{cases} z(\delta^2 - z^2)^2 & , |z| < \delta \\ 0 & , |z| \geq \delta \end{cases}$$

Non-convex, non-diff., but robust to outliers.

**Hinge loss:**

$$[1 - y_n f(\mathbf{x}_n)]_+ = \max(0, 1 - y_n f(\mathbf{x}_n))$$

**Logistic loss:**  $\log(1 - \exp(y_n f(\mathbf{x}_n)))$

### 3 Linear Regression

- Model that assume linear relationship

$$\mathbf{y}_n \approx f(\mathbf{x}_n) := \mathbf{w}_0 + \mathbf{w}_1 \mathbf{x}_{n1} + \dots = \mathbf{w}_0 + \mathbf{x}_n^T \mathbf{w}$$

$$\approx \tilde{\mathbf{x}}_n^T \mathbf{w}, \text{ where } \tilde{\mathbf{x}} \text{ contains offset comp.}$$

- $D > N$  problem: task is underdetermined.

### 4 Optimization

- Local minimum:  
 $L(w^*) \leq L(w) \forall w : \|w - w^*\| < \epsilon$
- Global minimum:  $L(w^*) \leq L(w) \forall w$

### 4.1 Grid search

- Compute the cost over a grid of  $V$  points.  
Exponential Complexity  $\mathcal{O}(|V|^D)$ ,  $D$  is the dimension. Hard to find a good range of values. No guarantee to converge.

### 4.2 GD - Gradient Descent (Batch)

- GD uses only first-order information and takes steps in the opposite direction of the gradient
- Given cost function  $\mathcal{L}(\mathbf{w})$  we want to find  $\mathbf{w} = \arg \min_{\mathbf{w}} \mathcal{L}(\mathbf{w})$
- Take steps in the opposite direction of the gradient

$$\mathbf{w}^{(t+1)} \leftarrow \mathbf{w}^{(t)} - \gamma \nabla \mathcal{L}(\mathbf{w}^{(t)})$$

- With  $\gamma$  too big, method might diverge. With  $\gamma$  too small, convergence is slow.
- Very sensitive to ill-conditioning  $\Rightarrow$  always normalize features  $\Rightarrow$  allow different directions to converge at same speed.

### 4.3 SGD - Stochastic Gradient Descent

SGD update rule (only n-th training example):

$$\mathbf{w}^{(t+1)} \leftarrow \mathbf{w}^{(t)} - \gamma \nabla \mathcal{L}_n(\mathbf{w}^{(t)})$$

*Idea:* Cheap but unbiased estimate of grad.

$$\mathbb{E}[\nabla \mathcal{L}_n(\mathbf{w})] = \nabla(\mathbf{w})$$

Robbins-Monroe condition:

- $\gamma^{(t)} : \sum_{t=1}^{\infty} \gamma^{(t)} = \infty; \sum_{t=1}^{\infty} (\gamma^{(t)})^2 < \infty$
- e.g.  $\gamma^{(t)} = 1/(t+1)^r, r \in (0.5, 1)$

### 4.4 Mini-batch SGD

Update direction ( $B \subseteq [N]$ ):

$$\mathbf{g}^{(t)} := \frac{1}{|B|} \sum_{n \in B} \nabla \mathcal{L}_n(\mathbf{w}^{(t)})$$

Update rule :  $\mathbf{w}^{(t+1)} \leftarrow \mathbf{w}^{(t)} - \gamma \mathbf{g}^{(t)}$

### 4.5 Gradients for MSE

- We define the error vector  $\mathbf{e}$ :

$$\mathbf{e} := \mathbf{y} - \mathbf{X} \mathbf{w}$$

- and MSE as follows:

$$\mathcal{L}(\mathbf{w}) = \frac{1}{2N} \sum_{n=1}^N (\mathbf{y}_n - \tilde{\mathbf{x}}_n^T \mathbf{w})^2 = \frac{1}{2N} \mathbf{e}^T \mathbf{e}$$

- then the gradient is given by

$$\nabla \mathcal{L}(\mathbf{w}) = -\frac{1}{N} \mathbf{X}^T \mathbf{e}$$

- Optimality conditions:

1. *necessary*: gradient equal zero:

$$\frac{d\mathcal{L}(\mathbf{w}^*)}{d\mathbf{w}} = 0$$

2. *sufficient*: Hessian matrix is positive

$$\text{definite: } \mathbf{H}(\mathbf{w}^*) = \frac{d^2 \mathcal{L}(\mathbf{w}^*)}{d\mathbf{w} d\mathbf{w}^T} = \frac{1}{N} \mathbf{X}^T \mathbf{X}$$

## 4.6 Subgradients (Non-Smooth OPT)

A vector  $\mathbf{g} \in \mathbb{R}^D$  s.t.

$$\mathcal{L}(\mathbf{u}) \geq \mathcal{L}(\mathbf{w}) + \mathbf{g}^T (\mathbf{u} - \mathbf{w}) \quad \forall \mathbf{u} \in \mathbb{R}^D$$

is the subgradient to  $\mathcal{L}$  at  $\mathbf{w}$ . If  $\mathcal{L}$  is differentiable at  $\mathbf{w}$ , we have  $\mathbf{g} = \nabla \mathcal{L}(\mathbf{w})$

### 4.7 Constrained Optimization

Find solution  $\min \mathcal{L}(\mathbf{w})$  s.t.  $\mathbf{w} \in \mathcal{C}$

- Add proj. onto  $\mathcal{C}$  after each step:

$$P_{\mathcal{C}}(\mathbf{w}') = \arg \min_{\mathbf{v}} \|\mathbf{v} - \mathbf{w}'\|, \mathbf{v} \in \mathcal{C}$$

$$\mathbf{w}^{(t+1)} = P_{\mathcal{C}}[\mathbf{w}^{(t)} - \gamma \nabla \mathcal{L}(\mathbf{w}^{(t)})]$$

- Use penalty functions

- $\min \mathcal{L}(\mathbf{w}) + I_{\mathcal{C}}, I_{\mathcal{C}} = 0$  if  $\mathbf{w} \in \mathcal{C}$ ,  $\infty$  otherwise
- $\min \mathcal{L}(\mathbf{w}) + \lambda \|\mathbf{Aw} - \mathbf{b}\|$
- Stopping criteria when  $\mathcal{L}(\mathbf{w})$  close to 0

### 4.8 Complexities for MSE/MAE per iteration

- GD= $\mathcal{O}(ND)$
- MB-GD= $\mathcal{O}(BD)$
- SGD= $\mathcal{O}(D)$

## 5 Least Squares

- Use the first optimality conditions:

$$\nabla \mathcal{L}(\mathbf{w}^*) = 0 \Rightarrow \mathbf{X}^T \mathbf{e} = \mathbf{X}^T (\mathbf{y} - \mathbf{X} \mathbf{w}) = 0$$

- When  $\mathbf{X}^T \mathbf{X}$  is invertible, we have the closed-form expression

$$\mathbf{w}^* = (\mathbf{X}^T \mathbf{X})^{-1} \mathbf{X}^T \mathbf{y}$$

- thus we can predict values for a new  $\mathbf{x}_m$

$$\mathbf{y}_m := \mathbf{x}_m^T \mathbf{w}^* = \mathbf{x}_m^T (\mathbf{X}^T \mathbf{X})^{-1} \mathbf{X}^T \mathbf{y}$$

- The **Gram matrix**  $\mathbf{X}^T \mathbf{X}$  is pd and is also invertible iff  $\mathbf{X}$  has full column rank.
- *Complexity*:  $\mathcal{O}(ND^2 + D^3) \equiv \mathcal{O}(ND^2)$
- $\mathbf{X}$  can be rank deficient when  $D > N$  or when the columns  $\tilde{\mathbf{x}}_d$  are nearly collinear.  $\Rightarrow$  matrix is ill-conditioned.
- Can still solve using a linear system solver using normal equations:

$$\mathbf{X}^T \mathbf{X} \mathbf{w} = \mathbf{X}^T \mathbf{y}$$

### 6 Maximum Likelihood (MLE)

- Let define the noise  $\epsilon_n \sim \mathcal{N}(0, \sigma^2)$ .

$$\rightarrow \mathbf{y}_n = \mathbf{x}_n^T \mathbf{w} + \epsilon_n$$

- Another way of expressing this:

$$p(\mathbf{y}|\mathbf{X}, \mathbf{w}) = \prod_{n=1}^N p(\mathbf{y}_n|\mathbf{x}_n, \mathbf{w})$$

$$= \prod_{n=1}^N \mathcal{N}(\mathbf{y}_n|\mathbf{x}_n^T \mathbf{w}, \sigma^2)$$

which defines the likelihood of observing  $\mathbf{y}$  given  $\mathbf{X}$  and  $\mathbf{w}$

- Define cost with log-likelihood

$$\mathcal{L}_{MLE}(\mathbf{w}) = \log p(\mathbf{y}|\mathbf{X}, \mathbf{w})$$

$$= -\frac{1}{2\sigma^2} \sum_{n=1}^N (\mathbf{y}_n - \mathbf{x}_n^T \mathbf{w})^2 + c n s t$$

- Maximum likelihood estimator (MLE) gives another way to design cost functions

$$\arg \min_{\mathbf{w}} \mathcal{L}_{MLE}(\mathbf{w}) = \arg \max_{\mathbf{w}} \mathcal{L}_{MLE}(\mathbf{w})$$

- MLE can also be interpreted as finding the model under which the observed data is most likely to have been generated from.
- $\mathbf{w}_{MLE} \rightarrow \mathbf{w}_{true}$  for large amount of data

## 7 Ridge Regression and LASSO

- Add **regularization term**

$$\min_{\mathbf{w}} \mathcal{L}(\mathbf{w}) + \mathbf{w}(\mathbf{w})$$

- This corresponds to MAP estimate with prior on weights.
- $L_2$ -Reg. (Ridge):  $\Omega(\mathbf{w}) = \lambda \|\mathbf{w}\|_2^2$
- $\rightarrow$  small values of  $\mathbf{w}_i$ , not sparse
- $\rightarrow \mathbf{w}^* = (\mathbf{X}^T \mathbf{X} + \lambda' \mathbf{I})^{-1} \mathbf{X}^T \mathbf{y}$  with  $\lambda' = 2N\lambda$

- $\rightarrow (\mathbf{X}^T \mathbf{X} + \lambda' \mathbf{I})^{-1}$  exists (lifted eigenvalues)
- $L_1$ -Reg. (Lasso):  $\Omega(\mathbf{w}) = \lambda \|\mathbf{w}\|_1$
- $\rightarrow$  sparsity of weight vector
- $\rightarrow$  implicit model selection
- **Maximum-a-posteriori (MAP)**
- (i) Posterior prob.  $\propto$  Likelihood  $\times$  Prior prob

$$p(\mathbf{y}|\mathbf{X} \mathbf{w}) = \prod_{n=1}^N \mathcal{N}(\mathbf{y}_n|\mathbf{x}_n^T \mathbf{w}, \sigma_n^2)$$

$$p(\mathbf{w}) = \mathcal{N}(\mathbf{w}|0, \sigma_0^2 \mathbf{I}_D)$$

$$\text{then } \rightarrow \mathbf{w}^* = \arg \max_{\mathbf{w}} p(\mathbf{y}|\mathbf{X} \mathbf{w}) \cdot p(\mathbf{w})$$

$$\mathbf{w}^* = \arg \min_{\mathbf{w}} \sum \frac{1}{2\sigma_n^2} (\mathbf{y}_n - \mathbf{x}_n^T \mathbf{w})^2 + \frac{1}{2\sigma_0^2} \|\mathbf{w}\|^2$$

## 8 Bias-Variance decomposition

- **Simple** (e.g. large  $\lambda$ )  $\rightarrow$  large bias but low variance
- **Complex** (e.g. small  $\lambda$ )  $\rightarrow$  low bias but large variance
- The expected squared loss between true model and learned model is a sum of three non-negative terms:

$$\mathbb{E}_S[(f(x) + \epsilon - f_S(x))^2] = \text{Var}[\epsilon] + \text{bias} + \text{variance:}$$

- **Bias**  $= (f(x) - \mathbb{E}_{S'}[f_{S'}(x)])^2$ : Difference between actual value and expected prediction.
- **Variance**  $= \mathbb{E}_S[(\mathbb{E}_{S'}[f_{S'}(x)] - f_S(x))^2]$ : variance of predictions between training sets.
- All terms are lower bounds for the error.
- Cannot do better than  $\text{Var}[\epsilon]$ .

## 9 Logistic Regression

- **Classification** relates input variables  $\mathbf{x}$  to discrete output variable  $\mathbf{y}$
- **Binary classifier**: we use  $y = 0$  for  $\mathbf{C}_1$  and  $y = 1$  for  $\mathbf{C}_2$ .
- Can use least-squares to predict  $\hat{y}_s$

$$\hat{y} = \begin{cases} \mathbf{C}_1 & \hat{y}_s < 0.5 \\ \mathbf{C}_2 & \hat{y}_s \geq 0.5 \end{cases}$$

- **Logistic function**

$$\sigma(x) = \frac{\exp(x)}{1 + \exp(x)}$$

$$p(\mathbf{y}_n = \mathbf{C}_1|\mathbf{x}_n) = \sigma(\mathbf{x}^T \mathbf{w})$$

- The probabilistic model:

$$p(\mathbf{y}|\mathbf{X}, \mathbf{w}) = \prod_{n=1}^N \sigma(\mathbf{x}_n^T \mathbf{w})^{y_n} (1 - \sigma(\mathbf{x}_n^T \mathbf{w}))^{1-y_n}$$

- The negative log-likelihood (w.r.t. MLE):

$$\mathcal{L}(\mathbf{w}) = - \sum_{n=1}^N \mathbf{y}_n \ln \sigma(\mathbf{x}_n^T \mathbf{w}) + (1 - \mathbf{y}_n) \ln(1 - \sigma(\mathbf{x}_n^T \mathbf{w}))$$

$$= - \sum_{n=1}^N \ln[1 + \exp(\mathbf{x}_n^T \mathbf{w})] - \mathbf{y}_n \mathbf{x}_n^T \mathbf{w}$$

- We can use the fact that

$$\frac{d}{dz} \ln(1 + \exp(z)) = \sigma(z)$$

- Gradient of the log-likelihood

$$\mathbf{g} = \nabla \mathcal{L}(\mathbf{w}) = \sum_{n=1}^N \mathbf{x}_n (\sigma(\mathbf{x}_n^T \mathbf{w}) - \mathbf{y}_n)$$

$$= \mathbf{X}^T [\sigma(\mathbf{X} \mathbf{w}) - \mathbf{y}]$$

- The neg. log-likelihood  $-\mathcal{L}_{MLE}(\mathbf{w})$  is convex
- **Hessian** of the neg. log-likelihood

- We know that
- $$\frac{d\sigma(t)}{dt} = \sigma(t)(1 - \sigma(t))$$

- Hessian is the derivative of the gradient

$$\mathbf{H}(\mathbf{w}) = \frac{d\mathbf{g}(\mathbf{w})}{d\mathbf{w}^T} = \sum_{n=1}^N \frac{d}{d\mathbf{w}^T} \mathbf{x}_n \sigma(\mathbf{x}_n^T \mathbf{w})$$

$$= \sum_{n=1}^N \mathbf{x}_n \mathbf{x}_n^T \sigma(\mathbf{x}_n^T \mathbf{w})(1 - \sigma(\mathbf{x}_n^T \mathbf{w}))$$

$$= \tilde{\mathbf{X}}^T \mathbf{S} \tilde{\mathbf{X}}$$

where  $\mathbf{S}$  is a  $N \times N$  diagonal matrix with diagonals

$$S_{nn} = \sigma(\mathbf{x}_n^T \mathbf{w})(1 - \sigma(\mathbf{x}_n^T \mathbf{w}))$$

- The neg. log-likelihood is not strictly convex. *????*

## Newton's Method

- Uses second-order information and takes steps in the direction that minimizes a quadratic approximation (Taylor)

## 11 k-Nearest Neighbor (k-NN)

- Performs best in low dimensions.
- Assumes close points have similar values
- The k-NN regressor:

$$f(\mathbf{x}) = \frac{1}{k} \sum_{\mathbf{x}_n \in nbh_k(\mathbf{x})} \mathbf{y}_n$$

- The k-NN classifier:

$$f(\mathbf{x}) = \text{modus } x_n | \mathbf{x}_n \in nbh_k(\mathbf{x})$$

- Large k  $\rightarrow$  smoothing over large area
- Small k  $\rightarrow$  averaging over small area
- **Curse of dimensionality:**

- As dimensionality grows fixed-size training set covers dwindling fraction of input space  $\Rightarrow$  If we want to consider fixed fraction  $\alpha$  of points and increase dimension, we need to explore almost whole range in each dimension.

- In high dimensions, points are far from each other  $\Rightarrow$  choice of NN becomes essentially random.

$$\text{Need radius} \\ r = \sqrt[d]{1 - \frac{1}{N^{\frac{1}{d}}}}$$

to have at least one data point in  $r^d$  rectangle with  $p \geq \frac{1}{2}$ .

- **NN performance:**

- Optimal classifier:  $f_*(x) = \mathbb{P}[y=1|x] > \frac{1}{2}$
- $\mathbb{E}_S[L(f_S)] \leq 2L(f_*) + 4c\sqrt{dN} \frac{1}{1+d}$

## 12 Support Vector Machine

- Assume  $y_n \in \{-1, 1\}$
- SVM optimizes the following cost

$$\mathcal{L}(\mathbf{w}) = \min_{\mathbf{w}} \sum_{n=1}^N [1 - \mathbf{y}_n \mathbf{x}_n^T \mathbf{w}]_+ + \frac{\lambda}{2} \|\mathbf{w}\|^2$$

- Can be optimised using subgradient descent.
- **Case: Linear separability:** We get a separating hyperplane, no point in the margin and  $w$ , s.t margin is maximised ( $2/\|w\|$ ).
- This is called hard-margin compared to soft-margin formulation.
- **Duality:**

- Hard to minimize  $g(\mathbf{w})$  so we define  $\mathcal{L}(\mathbf{w}) = \max_{\alpha} G(\mathbf{w}, \alpha)$

- we use the property that  $[\mathbf{v}_n]_+ = \max(0, \mathbf{v}_n) = \max_{\alpha_n \in [0, 1]} \alpha_n \mathbf{v}_n$

- We can rewrite the problem as

$$\min_{\mathbf{w}} \max_{\alpha} \sum_{n=1}^N \alpha_n (1 - \mathbf{y}_n \phi_n^T \mathbf{w}) + \frac{\lambda}{2} \|\mathbf{w}\|^2$$

- This is differentiable, convex in  $\mathbf{w}$  and concave in  $\alpha$

- **Minimax theorem:**

$$\min_{\mathbf{w}} \max_{\alpha} G(\mathbf{w}, \alpha) = \max_{\alpha} \min_{\mathbf{w}} G(\mathbf{w}, \alpha)$$

because  $G$  is convex in  $\mathbf{w}$  and concave in  $\alpha$ .

- Derivative w.r.t.  $\mathbf{w}$ :

$$\nabla_{\mathbf{w}} G(\mathbf{w}, \alpha) = - \sum_{n=1}^N \alpha_n \mathbf{y}_n \mathbf{x}_n + \lambda \mathbf{w}$$

- Equating this to 0, we get:

$$\mathbf{w}(\alpha) = \frac{1}{\lambda} \sum_{n=1}^N \alpha_n \mathbf{y}_n \mathbf{x}_n = \frac{1}{\lambda} \mathbf{X}^T \mathbf{Y} \alpha$$

$$\mathbf{Y} := \text{diag}(\mathbf{y})$$

- Plugging  $\mathbf{w}^*$  back in the dual problem

$$\max_{\alpha \in [0, 1]^N} \alpha^T \mathbf{1} - \frac{1}{2\lambda} \alpha^T \mathbf{Y} \mathbf{X} \mathbf{X}^T \mathbf{Y} \alpha$$

- Data only enters as  $\mathbf{K} = \mathbf{X}^T \mathbf{X}$ .

- **Non support vector:** Example that lies on the correct side, outside margin  $\alpha_n = 0$

- **Essen. support vector:** Example that lies on the margin  $\alpha_n \in (0, 1)$
- **Bound support vector:** Example that lies strictly inside the margin or wrong side  $\alpha_n = 1$
- Use Coordinates Descent to find  $\alpha$ . Update one coordinate (argmin) at the time and others constant.

## 13 Kernel Ridge Regression

- The following is true for ridge regression

$$\mathbf{w}^* = (\mathbf{X}^T \mathbf{X} + \lambda \mathbf{I}_D)^{-1} \mathbf{X}^T \mathbf{y}, \quad (1)$$

$$= \mathbf{X}^T (\mathbf{X} \mathbf{X}^T + \lambda \mathbf{I}_N)^{-1} \mathbf{y} = \mathbf{X}^T \alpha^*, \quad (2)$$

- Complexity of computing  $\mathbf{w}$ : (1)  $O(D^2 N + D^3)$ , (2)  $O(DN^2 + N^3)$
- Thus we have  $\mathbf{w}^* = \mathbf{X}^T \alpha^*$ , with  $\mathbf{w}^* \in \mathbb{R}^D$  and  $\alpha^* \in \mathbb{R}^N$
- The representer theorem allows us to write an equivalent optimization problem in terms of  $\alpha$ .

$$\alpha = \text{argmax}_{\alpha} \left( -\frac{1}{2} \alpha^T (\mathbf{X}^T \mathbf{X} + \lambda \mathbf{I}_N) \alpha + \alpha^T \mathbf{y} \right)$$

- $\mathbf{K} = \mathbf{X} \mathbf{X}^T$  is called the **kernel matrix** or **Gram matrix**.
- If  $\mathbf{K}$  is positive definite and symmetric, then it's called a **Mercer Kernel**.
- $\mathbf{K}_{i,j} = k(\mathbf{x}_i, \mathbf{x}_j)$
- If the kernel is Mercer, then there exists a function  $\phi(\mathbf{x})$  s.t.

$$k(\mathbf{x}, \mathbf{x}') = \phi(\mathbf{x})^T \phi(\mathbf{x}')$$

- **Kernel trick:**

- compute dot-product in  $\mathbb{R}^m$  while remaining in  $\mathbb{R}^n$
- Replace  $(\mathbf{x}, \mathbf{x}')$  with  $k(\mathbf{x}, \mathbf{x}')$ .
- **Common Kernel**

- $x \in \mathbb{R}, k(\mathbf{x}, \mathbf{x}') = (xx')^2 \Rightarrow \phi(x) = x^2$
- Radial Basis function kernel (RBF)

$$k(\mathbf{x}, \mathbf{x}') = \exp\left(-\frac{1}{2}(\mathbf{x} - \mathbf{x}')^T(\mathbf{x} - \mathbf{x}')\right)$$

- Thus we get

$$\mathbf{y} = \mathbf{w}^T \mathbf{x} = \sum_{i=1}^K \alpha_i \mathbf{x}_i^T \mathbf{x} = \sum_{i=1}^K \alpha_i k(\mathbf{x}, \mathbf{x}_i)$$

## 14 Unsupervised

Learn pattern without labels.

- **word2vec:** map every word to a vector  $w_i \in \mathbb{R}^K$ ,  $K$  large, that captures its semantics.

- **Topic model:** Documents consist of collections of topics

- topic = probability distribution over words
- use clustering to pick out representative topics

## 15 K-means

- **Unsupervised learning:** Represent particular input patterns in a way that reflects the statistical structure of the overall collections of input patterns.
- **Cluster** are groups of points whose inter-point distances are small compared to the distances outside the cluster.

$$\min_{\mathbf{z}, \mu} \mathcal{L}(\mathbf{z}, \mu) = \sum_{k=1}^K \sum_{n=1}^N z_{nk} \|\mathbf{x}_n - \mu_k\|_2^2$$

such that  $z_{nk} \in \{0, 1\}$  and  $\sum_{k=1}^K z_{nk} = 1$

- K-means algorithm (Coordinate Descent): Initialize  $\mu_k$ , then iterate

- For all  $n$ , compute  $\mathbf{z}_n$  given  $\mu$

$$z_{nk} = \begin{cases} 1 & \text{if } k = \text{argmin}_j \|\mathbf{x}_n - \mu_j\|_2^2 \\ 0 & \text{otherwise} \end{cases}$$

- For all  $k$ , compute  $\mu_k$  given  $\mathbf{z}$

$$\mu_k = \frac{\sum_{n=1}^N z_{nk} \mathbf{x}_n}{\sum_{n=1}^N z_{nk}}$$

- A good initialization procedure is to choose the prototypes to be equal to a random subset of  $K$  data points.
- Probabilistic model 
$$p(\mathbf{z}, \mu) = \prod_{n=1}^N \prod_{k=1}^K [\mathcal{N}(\mathbf{x}_n | \mu_k, \mathbf{I})]^{z_{nk}}$$
- $$-\log p(\mathbf{x}_n | \mu, \mathbf{z}) = \sum_k \sum_{k=1}^K \frac{1}{2} \|\mathbf{x}_n - \mu_k\|^2 z_{nk} + c'$$
- K-means as a Matrix Factorization 
$$\min_{\mathbf{z}, \mu} \mathcal{L}(\mathbf{z}, \mu) = \|\mathbf{X} - \mathbf{M} \mathbf{Z}^T\|_{\text{Frob}}^2$$
- Computation can be heavy, each example can belong to only on cluster and clusters have to be spherical.

## 16 Gaussian Mixture Models

- Clusters can be elliptical using a full covariance matrix instead of isotropic covariance.

$$p(\mathbf{X} | \mu, \Sigma, \mathbf{z}) = \prod_{n=1}^N \prod_{k=1}^K [\mathcal{N}(\mathbf{x}_n | \mu_k, \Sigma_k)]^{z_{nk}}$$

- **Soft-clustering:** Points can belong to several cluster by defining  $z_n$  to be a random variable.

$$p(z_n = k) = \pi_k \text{ where } \pi_k > 0, \forall k, \sum_{k=1}^K \pi_k = 1$$

- Joint distribution of Gaussian mixture model 
$$p(\mathbf{X}, \mathbf{z} | \mu, \Sigma, \pi) = \prod_{n=1}^N p(\mathbf{x}_n | \mathbf{r}_n, \mu, \Sigma) p(\mathbf{z}_n | \pi)$$

$$= \prod_{n=1}^N \prod_{k=1}^K [(\mathcal{N}(\mathbf{x}_n | \mu_k, \Sigma_k))^{z_{nk}}] \prod_{k=1}^K [\pi_k]^{z_{nk}}$$

- $z_n$  are called *latent* unobserved variables
- Unknown parameters are given by  $\theta = \{\mu, \Sigma, \pi\}$
- We get the **marginal likelihood** by marginalizing  $z_n$  out from the likelihood

$$\begin{aligned} p(\mathbf{x}_n | \theta) &= \sum_{k=1}^K p(\mathbf{x}_n, z_n = k | \theta) \\ &= \sum_{k=1}^K p(z_n = k | \theta) p(\mathbf{x}_n | z_n = k, \theta) \\ &= \sum_{k=1}^K \pi_k \mathcal{N}(\mathbf{x}_n | \mu_k, \Sigma_k) \end{aligned}$$

- Without a latent variable model, number of parameters grow at rate  $O(N)$
- After marginalization, the growth is reduced to  $O(D^2 K)$
- To get maximum likelihood estimate of  $\theta$ , we maximize

$$\max_{\theta} \sum_{n=1}^N \log \sum_{k=1}^K \pi_k \mathcal{N}(\mathbf{x}_n | \mu_k, \Sigma_k)$$

## 17 Expectation Maximization Algorithm

- [ALGORITHM] Start with  $\theta^{(1)}$  and iterate

- Expectation step:* Compute a lower bound to the cost such that it is tight at the previous  $\theta^{(t)}$  with equality when,

$$q_{kn} = \frac{\pi_k \mathcal{N}(\mathbf{x}_n | \mu_k, \Sigma_k)}{\sum_{k=1}^K \pi_k \mathcal{N}(\mathbf{x}_n | \mu_k, \Sigma_k)}$$

- Maximization step:* Update  $\theta$

$$\theta^{(t+1)} = \text{argmax}_{\theta} \mathcal{L}(\theta, \theta^{(t)})$$

$$\mu_k^{(t+1)} = \frac{\sum_{n=1}^N q_{kn}^{(t)} (\mathbf{x}_n) \mathbf{x}_n}{\sum_{n=1}^N q_{kn}^{(t)}}$$

$$\Sigma_k^{(t+1)} = \frac{\sum_{n=1}^N q_{kn}^{(t)} (\mathbf{x}_n - \mu_k^{(t+1)}) (\mathbf{x}_n - \mu_k^{(t+1)})^T}{\sum_{n=1}^N q_{kn}^{(t)}}$$

$$\pi_k^{(t+1)} = \frac{1}{N} \sum_{n=1}^N q_{kn}^{(t)}$$

- If covariance is diagonal  $\rightarrow$  K-means.

## 18 Matrix factorization

- We have  $D$  movies and  $N$  users
- $\mathbf{X}$  is a matrix  $D \times N$  with  $x_{dn}$  the rating of  $n$ 'th user for  $d$ 'th movie.
- We project data vectors  $\mathbf{x}_n$  to a smaller dimension  $\mathbf{z}_n \in \mathbb{R}^M$
- We have now 2 latent variables:
  - $\mathbf{Z}$  a  $N \times K$  matrix that gives features for the users
  - $\mathbf{W}$  a  $D \times K$  matrix that gives features for the movies

- We can add a regularizer and minimize the following cost:

$$\begin{aligned} \mathcal{L}(\mathbf{W}, \mathbf{Z}) &= \frac{1}{2} \sum_{(d,n) \in \Omega} [x_{dn} - (\mathbf{W} \mathbf{Z}^T)_{dn}]^2 \\ &\quad + \frac{\lambda_w}{2} \|\mathbf{W}\|_{\text{Frob}}^2 + \frac{\lambda_z}{2} \|\mathbf{Z}\|_{\text{Frob}}^2 \end{aligned}$$

- **SGD:** For one fixed element  $(d, n)$  we derive entry  $(d', k)$  of  $\mathbf{W}$  (if  $d = d'$  oth. 0):

$$\begin{aligned} \frac{\partial}{\partial w_{d',k}} f_{d,n}(\mathbf{W}, \mathbf{Z}) &= -[x_{dn} - (\mathbf{W} \mathbf{Z}^T)_{dn}] z_{nk} \\ \text{And of } \mathbf{Z} \text{ (if } n = n' \text{ oth. 0):} \\ \frac{\partial}{\partial z_{n',k}} f_{d,n}(\mathbf{W}, \mathbf{Z}) &= -[x_{dn} - (\mathbf{W} \mathbf{Z}^T)_{dn}] w_{nk} \end{aligned}$$

$$\mathbf{W}^{t+1} = \mathbf{W}^t - \gamma \nabla_w f_{d,n}(\mathbf{W}^t, \mathbf{Z}^t)$$

$$\mathbf{Z}^{t+1} = \mathbf{Z}^t - \gamma \nabla_z f_{d,n}(\mathbf{W}^t, \mathbf{Z}^t)$$

- We can use coordinate descent algorithm, by first minimizing w.r.t.  $\mathbf{Z}$  given  $\mathbf{W}$  and then minimizing  $\mathbf{W}$  given  $\mathbf{Z}$ . This is called

**Alternating least-squares (ALS):**

$$\mathbf{Z}^T \leftarrow (\mathbf{W}^T \mathbf{W} + \lambda_z \mathbf{I}_K)^{-1} \mathbf{W}^T \mathbf{X}$$

$$\mathbf{W}^T \leftarrow (\mathbf{Z}^T \mathbf{Z} + \lambda_w \mathbf{I}_K)^{-1} \mathbf{Z}^T \mathbf{X}^T$$

- *Complexity:*  $O(DN K^2 + N K^3) \rightarrow O(DN K^2)$

## 19 Singular Value Decomposition

- Matrix factorization method  $\mathbf{X} = \mathbf{U} \mathbf{S} \mathbf{V}^T$

- $\mathbf{U}$  is a unitary  $D \times D$  matrix
- $\mathbf{V}$  is a unitary  $N \times N$  matrix
- $\mathbf{S}$  is a non-negative diagonal matrix of size  $D \times N$  which are called **singular values** appearing in a descending order.
- Columns of  $\mathbf{U}$  and  $\mathbf{V}$  are the left and right **singular vectors** respectively.

- Assuming  $D < N$  we have

$$\mathbf{X} = \sum_{d=1}^D s_d \mathbf{u}_d \mathbf{v}_d^T$$

- This tells you about the spectrum of  $\mathbf{X}$  where higher singular vectors contain the *low-frequency information* and lower singular values contain the *high-frequency information*.

- **Truncated SVD:**

Take the matrix  $\mathbf{S}^{(K)}$  with the  $K$  first diagonal elements non zero. Then, rank- $K$  approx:

$$\mathbf{X} \approx \mathbf{X}_K = \mathbf{U} \mathbf{S}^{(K)} \mathbf{V}^T$$

## 20 Principal Component Analysis

- PCA is a dimensionality reduction method and a method to decorrelate the data
- $\mathbf{X} \approx \tilde{\mathbf{X}} = \mathbf{W} \mathbf{Z}^T$  such that columns of  $\mathbf{W}$  are orthogonal.

- If the data is zero mean

$$\Sigma = \frac{1}{N} \mathbf{X} \mathbf{X}^T \Rightarrow \mathbf{X} \mathbf{X}^T = \mathbf{U} \mathbf{S}^2 \mathbf{U}^T$$

$$\Rightarrow \mathbf{U}^T \mathbf{X} \mathbf{X}^T \mathbf{U} = \mathbf{U}^T \mathbf{U} \mathbf{S}^2 \mathbf{U}^T \mathbf{U} = \mathbf{S}^2$$

- Thus the columns of matrix  $\mathbf{U}$  are called the **principal components** and they decorrelate the covariance matrix.
- Using SVD, we can compute the matrices in the following way

$$\mathbf{W} = \mathbf{U} \mathbf{S}_D^{1/2}, \mathbf{Z}^T = \mathbf{S}^{1/2} \mathbf{V}^T$$

- Not invariant under scalings of the feature = arbitrariness,  $\rightarrow$  normalize  $\mathbf{X}$

## 21 Neural Net

- Basic structure: One *input* layer of size  $D$ ,  $L$  *hidden* layers of size  $K$ , and one *output* layer. (*feedforward* network).
- $x_j^{(l)} = \phi \left( \sum_i w_{ij}^{(l)} x_i^{(l-1)} + b_j^{(l)} \right)$ .
- NN can represent the Riemmann sum with only two layers  $\Rightarrow$  It's powerful!
- Cost function: 
$$\mathcal{L} = \frac{1}{N} \sum_{n=1}^N \left( y_n - f^{(L+1)} \circ \dots \circ f^{(1)}(\mathbf{x}_n^{(0)}) \right)^2$$
 We can use SGD to minimize the cost function.

### 21.1 Backpropagation Algorithm

- *Forward pass:* Compute  $\mathbf{z}^{(l)} = (\mathbf{W}^{(l)})^T \mathbf{x}^{(l-1)} + \mathbf{b}^{(l)}$  with  $\mathbf{x}^{(0)} = \mathbf{x}_n$  and  $\mathbf{x}^{(l)} = \phi(\mathbf{z}^{(l)})$ .
- *Backward pass:* Set  $\delta^{(L+1)} = -2(y_n - \mathbf{x}^{(L+1)}) \phi'(z^{(L+1)})$  (if squared loss). Then compute

$$\begin{aligned} \delta_j^{(l)} &= \frac{\partial \mathcal{L}_n}{\partial z_j^{(l)}} = \sum_k \frac{\partial \mathcal{L}_n}{\partial z_k^{(l+1)}} \frac{\partial z_k^{(l+1)}}{\partial z_j^{(l)}} \\ &= \sum_k \delta_k^{(l+1)} \mathbf{w}_{j,k}^{(l+1)} \phi'(z_j^{(l)}) \end{aligned}$$

- *Final Computation:*

$$\frac{\partial \mathcal{L}_n}{\partial w_{i,j}^{(l)}} = \sum_k \frac{\partial \mathcal{L}_n}{\partial z_k^{(l)}} \frac{\partial z_k^{(l)}}{\partial w_{i,j}^{(l)}} = \frac{\partial \mathcal{L}_n}{\partial z_j^{(l)}} \frac{\partial z_j^{(l)}}{\partial w_{i,j}^{(l)}} = \delta_j^{(l)} \mathbf{x}_i^{(l-1)}$$

$$\frac{\partial \mathcal{L}_n}{\partial b_j^{(l)}} = \sum_k \frac{\partial \mathcal{L}_n}{\partial z_k^{(l)}} \frac{\partial z_k^{(l)}}{\partial b_j^{(l)}} = \frac{\partial \mathcal{L}_n}{\partial z_j^{(l)}} \frac{\partial z_j^{(l)}}{\partial b_j^{(l)}} = \delta_j^{(l)} \cdot 1 = \delta_j^{(l)}$$

### 21.2 Activation Functions

**Sigmoid**  $\phi(x) = \frac{1}{1+e^{-x}}$  Positive, bounded.

$\phi'(x) \simeq 0$  for large  $|x| \Rightarrow$  Learning slow.

**Tanh**  $\tanh(x) = \frac{e^x - e^{-x}}{e^x + e^{-x}} = \phi(2x) - 1/2$ .

Balanced, bounded. Learning slow too.

**ReLU**  $(\mathbf{x})_+ = \max(0, x)$  Positive, unbounded. Derivate = 1 if  $x > 0$ , 0 if  $x < 0$

**Leaky ReLU**  $f(x) = \max(\alpha x, x)$  Remove 0 derivative.

**Maxout**

$f(\mathbf{x}) = \max_{\alpha} \mathbf{x}^T \mathbf{w}_1 + b_1, \dots, \mathbf{x}^T \mathbf{w}_k + b_k$  (Generalization of ReLU)

### 21.3 Convolutional NN

Sparse connections and *weights sharing*: reduce complexity. (e.g. pixels in pictures only depend on neighbours)

### 21.4 Reg, Data Augmentation and Dropout

- Regularization term:  $\frac{1}{2} \sum_{l=1}^{L+1} \mu^{(l)} \|\mathbf{W}^{(l)}\|_F^2$
- Weight decay is  $\Theta[t](1 - \eta \mu)$  in:  $\Theta[t+1] = \Theta[t] + \eta(\nabla \mathcal{L} + \mu \Theta[t])$
- Data Augm.: e.g. shift or rotation of pics

- Dropout: avoid overfit. Drop nodes randomly. (Then average multiple drop-NN)

**22 Bayes Net**

- Graph example:  $p(x, y, z) = p(y|x)p(z|x)p(x)$   
:  $(y \leftarrow x \rightarrow z)$
- **D-Separation** X and Y are D-separated by Z if every path from  $x \in X$  to  $y \in Y$  is blocked by Z. ( $\rightarrow$  independent)
- **Blocked Path** contains a variable that
  - is in Z and is **head-to-tail** or **tail-to-tail**
  - the node is **head-to-head** and neither the node nor any of its descendants are in Z.
- **Markov Blanket** (which blocks node A from the rest of the net) contains:
  - parents of A
  - children of A
  - parents of children of A