

# PCML Cheat Sheet

## 1 Math Prerequisites

- Bayes rule

$$p(A, B) = \underbrace{p(A|B)}_{\text{Lik.}} \underbrace{p(B)}_{\text{Prior}} = \underbrace{p(B|A)}_{\text{Post}} \underbrace{p(A)}_{\text{Marg. Lik.}}$$

- Gaussian distribution

$$\mathcal{N}(x|\mu, \sigma^2) = \frac{1}{\sqrt{2\pi\sigma^2}} \exp\left(-\frac{(x-\mu)^2}{2\sigma^2}\right)$$

$$\mathcal{N}(\mathbf{x}|\boldsymbol{\mu}, \boldsymbol{\Sigma}) = \frac{1}{\sqrt{2\pi|\boldsymbol{\Sigma}|}} \exp\left(-\frac{1}{2}(\mathbf{x}-\boldsymbol{\mu})^T \boldsymbol{\Sigma}^{-1}(\mathbf{x}-\boldsymbol{\mu})\right)$$

- Production of independent variables:  
 $\text{Var}(XY) = \mathbb{E}(X^2) \mathbb{E}(Y^2) - [\mathbb{E}(X)]^2 [\mathbb{E}(Y)]^2$
- Covariance matrix of a data vector  $\mathbf{x}$

$$\boldsymbol{\Sigma} = \frac{1}{N} \sum_{n=1}^N (\mathbf{x}_n - \mathbb{E}(\mathbf{x}))(\mathbf{x}_n - \mathbb{E}(\mathbf{x}))^T$$

### 1.1 Convexity

- A function  $f(x)$  is convex, if for any  $x_1, x_2 \in \mathbf{X}$  and for any  $0 \leq \lambda \leq 1$ , we have :

$$f(\lambda x_1 + (1-\lambda)x_2) \leq \lambda f(x_1) + (1-\lambda)f(x_2)$$

- The Hessian of a convex function is psd and for a strictly-convex function it's pd.

$$\mathbf{H}_{i,j} = d^2 f / dx_i dx_j$$

### 1.2 Linear Algebra

- **Condition number** If  $\mathbf{A}$  is normal ( $A^T A = A A^T$ ) then

$$k(\mathbf{A}) = \left| \frac{\lambda_{\max}(\mathbf{A})}{\lambda_{\min}(\mathbf{A})} \right|$$

- A positive definite matrix is **symmetric** with all positive eigenvalues
- The real symmetric  $N \times N$  matrix  $\mathbf{V}$  is said to be **positive semidefinite** if

$$\mathbf{a}^T \mathbf{V} \mathbf{a} \geq 0$$

- for any real  $N \times 1$  vector  $\mathbf{a}$ .
- **positive definite** if  $\mathbf{a}^T \mathbf{V} \mathbf{a} > 0$

### 2 Cost functions

- Cost functions are used to learn parameters that explain the data well.
- It is essential to make sure that a global minimum exist  $\rightarrow$  lower bounded

#### Mean square error (MSE):

$$MSE(\mathbf{w}) = \sum_{n=1}^N (y_n - f(\mathbf{x}_n))^2$$

- MSE is **convex** thus it has only one global minimum value.
- MSE is not good when outliers are present.

#### Mean Absolute Error (MAE):

$$MAE = \sum_{n=1}^N |y_n - f(\mathbf{x}_n)|$$

#### Huber loss

$$Huber = \begin{cases} \frac{1}{2} z^2 & , |z| \leq \delta \\ \delta |z| - \frac{1}{2} \delta^2 & , |z| > \delta \end{cases}$$

- Huber loss is convex, differentiable, and also robust to outliers but hard to set  $\delta$ .

#### Tukey's bisquare loss

$$L(z) = \begin{cases} z(\delta^2 - z^2)^2 & , |z| < \delta \\ 0 & , |z| \geq \delta \end{cases}$$

Non-convex, non-diff., but robust to outliers.

#### Hinge loss

$$Hinge = [1 - y_n f(\mathbf{x}_n)]_+ = \max(0, 1 - y_n f(\mathbf{x}_n))$$

#### Logistic loss

$$Logistic = \log(1 - \exp(y_n f(\mathbf{x}_n)))$$

## 3 Regression

- **Data** consists of  $N$  pairs  $(y_n, \mathbf{x}_n)$

1.  $y_n$  the  $n$ 'th output
  2.  $\mathbf{x}_n$  is a vector of  $D$  inputs
- **Prediction**: predict the output for a new input vector.
  - **Interpretation**: understand the effect of inputs on output.
  - **Outliers** are data that are far away from most of the other examples.

### 3.1 Linear Regression

- Model that assume linear relationship between inputs and the output.
- $y_n \approx f(\mathbf{x}_n) := w_0 + w_1 x_{n1} + \dots = w_0 + \mathbf{x}_n^T \mathbf{w}$  with  $\mathbf{w}$  the parameters of the model.
- Variance grows only linearly with dimensionality

### 4 Optimization

#### 4.1 Grid search

- Compute the cost over a grid of  $M$  points to find the minimum. Exponential Complexity. Hard to find a good range of values

#### 4.2 Gradient Descent

- GD uses only first-order information and takes steps in the opposite direction of the gradient
- Given cost function  $\mathcal{L}(\mathbf{w})$  we want to find  $\mathbf{w} = \arg \min_{\mathbf{w}} \mathcal{L}(\mathbf{w})$

#### 4.3 Batch Gradient Descent

- Take steps in the opposite direction of the gradient

$$\mathbf{w}^{(t+1)} \leftarrow \mathbf{w}^{(t)} - \gamma \nabla \mathcal{L}(\mathbf{w}^{(t)})$$

- with  $\gamma > 0$  the learning rate.
- With  $\gamma$  too big, method might diverge. With  $\gamma$  too small, convergence is slow.

#### 4.4 Gradients for MSE

- We define the error vector  $\mathbf{e}$ :  
 $\mathbf{e} := \mathbf{y} - \mathbf{X}\mathbf{w}$
- and MSE as follows:

$$\mathcal{L}(\mathbf{w}) = \frac{1}{2N} \sum_{n=1}^N (y_n - \tilde{\mathbf{x}}_n^T \mathbf{w})^2 = \frac{1}{2N} \mathbf{e}^T \mathbf{e}$$

- then the gradient is given by

$$\nabla \mathcal{L}(\mathbf{w}) = -\frac{1}{N} \mathbf{X}^T \mathbf{e}$$

- Optimality conditions:

1. *necessary*: gradient equal zero:  
 $\frac{d\mathcal{L}(\mathbf{w}^*)}{d\mathbf{w}} = 0$
2. *sufficient*: Hessian matrix is positive definite:  $\mathbf{H}(\mathbf{w}^*) = \frac{d^2 \mathcal{L}(\mathbf{w}^*)}{d\mathbf{w} d\mathbf{w}^T}$
- Very sensitive to illconditioning  $\Rightarrow$  always normalize features.
- *Complexity*:  $O(NDI)$  with  $I$  the number of iterations

### 4.5 Stochastic Gradient Descent

In ML, most cost functions are formulated as a sum over the training examples:

$$\mathcal{L}(\mathbf{w}) = \frac{1}{N} \sum_{n=1}^N \mathcal{L}_n(\mathbf{w})$$

$\Rightarrow$  SGD algo is given by update rule:

$$\mathbf{w}^{(t+1)} \leftarrow \mathbf{w}^{(t)} - \gamma \nabla \mathcal{L}_n(\mathbf{w}^{(t)})$$

Idea: Cheap but unbiased estimate of grad.

$$\mathbb{E}[\nabla \mathcal{L}_n(\mathbf{w})] = \nabla \mathcal{L}(\mathbf{w})$$

## 4.6 Mini-batch SGD

Update direction ( $B \subseteq [N]$ ):

$$\mathbf{g}^{(t)} := \frac{1}{|B|} \sum_{n \in B} \nabla \mathcal{L}_n(\mathbf{w}^{(t)})$$

Update rule :  $\mathbf{w}^{(t+1)} \leftarrow \mathbf{w}^{(t)} - \gamma \mathbf{g}^{(t)}$

### 4.7 Subgradients (Non-Smooth OPT)

A vector  $\mathbf{g} \in \mathbb{R}^D$  s.t.

$$\mathcal{L}(\mathbf{u}) \geq \mathcal{L}(\mathbf{w}) + \mathbf{g}^T (\mathbf{u} - \mathbf{w}) \quad \forall \mathbf{u}$$

is the subgradient to  $\mathcal{L}$  at  $\mathbf{w}$ . If  $\mathcal{L}$  is differentiable at  $\mathbf{w}$ , we have  $\mathbf{g} = \nabla \mathcal{L}(\mathbf{w})$

## 5 Least Squares

- Use the first optimality conditions:  
 $\nabla \mathcal{L}(\mathbf{w}^*) = 0 \Rightarrow \mathbf{X}^T \mathbf{e} = \mathbf{X}^T (\mathbf{y} - \mathbf{X}\mathbf{w}) = 0$
- When  $\mathbf{X}^T \mathbf{X}$  is invertible, we have the closed-form expression

$$\mathbf{w}^* = (\mathbf{X}^T \mathbf{X})^{-1} \mathbf{X}^T \mathbf{y}$$

- thus we can predict values for a new  $\mathbf{x}_m$   
 $y_m := \mathbf{x}_m^T \mathbf{w}^* = \mathbf{x}_m^T (\mathbf{X}^T \mathbf{X})^{-1} \mathbf{X}^T \mathbf{y}$
- The **Gram matrix**  $\mathbf{X}^T \mathbf{X}$  is pd and is also invertible iff  $\mathbf{X}$  has full column rank.
- *Complexity*:  $O(ND^2 + D^3) \equiv O(ND^2)$
- $\mathbf{X}$  can be rank deficient when  $D > N$  or when the columns  $\tilde{\mathbf{x}}_d$  are nearly collinear.  $\Rightarrow$  matrix is ill-conditioned.

## 6 Maximum Likelihood

- Let define our mistakes  $\epsilon_n \sim \mathcal{N}(0, \sigma^2)$ .  
 $\rightarrow y_n = \mathbf{x}_n^T \mathbf{w} + \epsilon_n$
- Another way of expressing this:

$$p(\mathbf{y}|\mathbf{X}, \mathbf{w}) = \prod_{n=1}^N p(y_n|\mathbf{x}_n, \mathbf{w}) = \prod_{n=1}^N \mathcal{N}(y_n|\mathbf{x}_n^T \mathbf{w}, \sigma^2)$$

which defines the likelihood of observing  $\mathbf{y}$  given  $\mathbf{X}$  and  $\mathbf{w}$

- Define cost with log-likelihood

$$\mathcal{L}_{lik}(\mathbf{w}) = \log p(\mathbf{y}|\mathbf{X}, \mathbf{w})$$

$$= -\frac{1}{2\sigma^2} \sum_{n=1}^N (y_n - \mathbf{x}_n^T \mathbf{w})^2 + c_{nst}$$

- Maximum likelihood estimator (MLE) gives another way to design cost functions  
 $\arg \min_{\mathbf{w}} \mathcal{L}_{MSE}(\mathbf{w}) = \arg \min_{\mathbf{w}} \mathcal{L}_{lik}(\mathbf{w})$
- MLE can also be interpreted as finding the model under which the observed data is most likely to have been generated from.
- With Laplace distribution

$$p(y_n|\mathbf{x}_n, \mathbf{w}) = \frac{1}{2b} e^{-\frac{1}{b}|y_n - \mathbf{x}_n^T \mathbf{w}|}$$

$$\sum_n \log p(y_n|\mathbf{x}_n, \mathbf{w}) = \sum_n |y_n - \mathbf{x}_n^T \mathbf{w}| + c_{nst}$$

## 7 Ridge Regression

- Linear models usually overfit. One way is to use nonlinear basis functions instead.

$$y_n = w_0 + \sum_{j=1}^M w_j \phi_j(\mathbf{x}_n) = \tilde{\boldsymbol{\phi}}(\mathbf{x}_n)^T \mathbf{w}$$

- This model is linear in  $\mathbf{w}$  but nonlinear in  $\mathbf{x}$ . Dimension is now  $M$ , not  $N$ .
- Polynomial basis  
 $\boldsymbol{\phi}(x_n) = [1, x_n, x_n^2, \dots, x_n^M]$
- The least square solution becomes  
 $\mathbf{w}_{lse}^* = (\tilde{\boldsymbol{\Phi}}^T \tilde{\boldsymbol{\Phi}})^{-1} \tilde{\boldsymbol{\Phi}}^T \mathbf{y}$

- Complex models overfit easily. Thus we can penalize them with a **regularization term**

$$\min_{\mathbf{w}} \left( \mathcal{L}(\mathbf{w}) + \frac{\lambda}{2N} \sum_{j=1}^M w_j^2 \right)$$

$$\mathbf{w}^* = \arg \min_{\mathbf{w}} \left( \frac{1}{2} (\mathbf{y} - \mathbf{X}\mathbf{w})^T (\mathbf{y} - \mathbf{X}\mathbf{w}) + \frac{\lambda}{2} \mathbf{w}^T \mathbf{w} \right)$$

- Note that  $w_0$  is not penalized.
- By differentiating and setting to zero we get

$$\mathbf{w}_{ridge} = (\tilde{\boldsymbol{\Phi}}^T \tilde{\boldsymbol{\Phi}} + \boldsymbol{\Lambda})^{-1} \tilde{\boldsymbol{\Phi}}^T \mathbf{y}$$

$$\boldsymbol{\Lambda} = \begin{bmatrix} 0 & \mathbf{0} \\ \mathbf{0} & \lambda I_M \end{bmatrix}$$

- Ridge regression improves the condition number of the Gram matrix since the eigenvalues of  $(\tilde{\boldsymbol{\Phi}}^T \tilde{\boldsymbol{\Phi}} + \lambda I_m)$  are at least  $\lambda$
- **Maximum-a-posteriori (MAP) estimator**:

- Maximizes the product of the likelihood and the **prior**.

$$\mathbf{w}_{MAP} = \arg \max_{\mathbf{w}} (p(\mathbf{y}|\mathbf{X}, \boldsymbol{\Lambda}) p(\mathbf{w}|\boldsymbol{\Sigma}))$$

- Assume  $w_0 = 0$

$$\mathbf{w}_{ridge} = \arg \max_{\mathbf{w}} \left( \log \left[ \prod_{n=1}^N \mathcal{N}(y_n|\mathbf{x}_n^T \mathbf{w}, \boldsymbol{\Lambda}) \times \mathcal{N}(\mathbf{w}|\mathbf{0}, \mathbf{I}) \right] \right)$$

- **Lasso regularizer** forces some  $w_i$  to be strictly 0 and therefore forces sparsity in the model.

$$\min_{\mathbf{w}} \frac{1}{2N} \sum_{n=1}^N (y_n - \tilde{\boldsymbol{\phi}}(\mathbf{x}_n)^T \mathbf{w})^2,$$

$$\text{such that } \sum_{i=1}^M |w_i| \leq \tau$$

## 8 Bias-Variance decomposition

- The expected test error can be expressed as the sum of two terms

- **Squared bias**: The average *shift* of the predictions
- **Variance**: measure how data points vary around their average.

$$\text{expected loss} = (\text{bias})^2 + \text{variance} + \text{noise}$$

- Model bias and estimation bias are important
- RR increases estimation bias and reduces var
- Model more complex increases test error
- Small  $\lambda \rightarrow$  low bias but large variance
- Large  $\lambda \rightarrow$  large bias but low variance
- $err = \sigma^2 + \mathbb{E}[f_{lse} - \mathbb{E}[f_{lse}]]^2 + [f_{true} - \mathbb{E}[f_{lse}]]^2$

## 9 Logistic Regression

- **Classification** relates input variables  $\mathbf{x}$  to discrete output variable  $y$
- **Binary classifier**: we use  $y = 0$  for  $\mathbf{C}_1$  and  $y = 1$  for  $\mathbf{C}_2$ .
- Can use least-squares to predict  $\hat{y}_*$

$$\hat{y} = \begin{cases} \mathbf{C}_1 & \hat{y}_* < 0.5 \\ \mathbf{C}_2 & \hat{y}_* \geq 0.5 \end{cases}$$

- **Logistic function**

$$\sigma(x) = \frac{\exp(x)}{1 + \exp(x)}$$

$$p(y_n = \mathbf{C}_1|\mathbf{x}_n) = \sigma(\mathbf{x}_n^T \mathbf{w})$$

$$p(y_n = \mathbf{C}_2|\mathbf{x}_n) = 1 - \sigma(\mathbf{x}_n^T \mathbf{w})$$

- The probabilistic model:

$$p(\mathbf{y}|\mathbf{X}, \mathbf{w}) = \prod_{n=1}^N \sigma(\mathbf{x}_n^T \mathbf{w})^{y_n} (1 - \sigma(\mathbf{x}_n^T \mathbf{w}))^{1-y_n}$$

- The log-likelihood:

$$\mathcal{L}_{MLE}(\mathbf{w}) = \sum_{n=1}^N (y_n \mathbf{x}_n^T \mathbf{w} - \log(1 + \exp(\mathbf{x}_n^T \mathbf{w})))$$

- We can use the fact that

$$\frac{d}{dx} \log(1 + \exp(x)) = \sigma(x)$$

- Gradient of the log-likelihood

$$\mathbf{g} = \frac{d\mathcal{L}}{d\mathbf{w}} = \sum_{n=1}^N (\mathbf{x}_n y_n - \mathbf{x}_n \sigma(\mathbf{x}_n^T \mathbf{w})) = -\mathbf{X}^T [\sigma(\mathbf{X}\mathbf{w}) - \mathbf{y}]$$

- The negative of the log-likelihood  $-\mathcal{L}_{MLE}(\mathbf{w})$  is convex

- **Hessian** of the log-likelihood

- We know that  
 $\frac{d\sigma(t)}{dt} = \sigma(t)(1 - \sigma(t))$
- Hessian is the derivative of the gradient

$$\mathbf{H}(\mathbf{w}) = -\frac{d\mathbf{g}(\mathbf{w})}{d\mathbf{w}^T} = \sum_{n=1}^N \frac{d}{d\mathbf{w}^T} \sigma(\mathbf{x}_n^T \mathbf{w}) \mathbf{x}_n$$

$$= \sum_{n=1}^N \mathbf{x}_n \sigma(\mathbf{x}_n^T \mathbf{w}) (1 - \sigma(\mathbf{x}_n^T \mathbf{w})) \mathbf{x}_n^T$$

$$= \tilde{\mathbf{X}}^T \mathbf{S} \tilde{\mathbf{X}}$$

where  $\mathbf{S}$  is a  $N \times N$  diagonal matrix with diagonals

- $S_{nn} = \sigma(\mathbf{x}_n^T \mathbf{w}) (1 - \sigma(\mathbf{x}_n^T \mathbf{w}))$
- The negative of the log-likelihood is not strictly convex.

- **Newton's Method**

- Uses second-order information and takes steps in the direction that minimizes a quadratic approximation

$$\mathcal{L}(\mathbf{w}) = \mathcal{L}(\mathbf{w}^{(k)}) + \nabla \mathcal{L}_k^T (\mathbf{w} - \mathbf{w}^{(k)})$$

$$+ (\mathbf{w} - \mathbf{w}^{(k)})^T \mathbf{H}_k (\mathbf{w} - \mathbf{w}^{(k)})$$

and it's minimum is at  $\mathbf{w}^{k+1} = \mathbf{w}^{(k)} - \gamma_k \mathbf{H}_k^{-1} \nabla \mathcal{L}_k$

- Complexity:  $O((ND^2 + D^3)I)$

- **Penalized Logistic Regression**

$$\min_{\mathbf{w}} \left( -\sum_{n=1}^N \log p(y_n|\mathbf{x}_n^T \mathbf{w}) + \lambda \sum_{d=1}^D w_d^2 \right)$$

## 10 Generalized Linear Model

### Exponential family distribution

$$p(\mathbf{y}|\boldsymbol{\eta}) = \frac{h(\mathbf{y})}{Z} \exp(\boldsymbol{\eta}^T \boldsymbol{\phi}(\mathbf{y}) - A(\boldsymbol{\eta}))$$

- Bernoulli distribution  
 $p(y|\mu) = \mu^y (1 - \mu)^{1-y}$

$$= \exp(y \log(\frac{\mu}{1-\mu}) + \log(1-\mu))$$

there is a relationship between  $\boldsymbol{\eta}$  and  $\boldsymbol{\mu}$  through the **link function**

$$\boldsymbol{\eta} = \log\left(\frac{\boldsymbol{\mu}}{1-\boldsymbol{\mu}}\right) \Leftrightarrow \boldsymbol{\mu} = \frac{e^{\boldsymbol{\eta}}}{1+e^{\boldsymbol{\eta}}}$$

- Note that  $\boldsymbol{\mu}$  is the mean parameter of  $y$

- Relationship between the mean  $\boldsymbol{\mu}$  and  $\boldsymbol{\eta}$  is defined using a link function  $g$

$$\boldsymbol{\eta} = \mathbf{g}(\boldsymbol{\mu}) \Leftrightarrow \boldsymbol{\mu} = \mathbf{g}^{-1}(\boldsymbol{\eta})$$

- First and second derivatives of  $A(\boldsymbol{\eta})$  are related to the mean and the variance

$$\frac{dA(\boldsymbol{\eta})}{d\boldsymbol{\eta}} = \mathbb{E}[\boldsymbol{\phi}(\boldsymbol{\eta})], \quad \frac{d^2 A(\boldsymbol{\eta})}{d\boldsymbol{\eta}^2} = \text{Var}[\boldsymbol{\phi}(\boldsymbol{\eta})]$$

- $A(\boldsymbol{\eta})$  is convex

- The generalized maximum likelihood cost to minimize is

$$\min_{\mathbf{w}} \mathcal{L}(\mathbf{w}) = -\sum_{n=1}^N \log(p(y_n|\mathbf{x}_n^T \mathbf{w}))$$

where  $p(y_n|\mathbf{x}_n^T \mathbf{w})$  is an exponential family distribution

We obtain the solution

$$\frac{d\mathcal{L}}{d\mathbf{w}} = \mathbf{X}^T [\mathbf{g}^{-1}(\boldsymbol{\eta}) - \boldsymbol{\phi}(\mathbf{y})]$$

11 **k-Nearest Neighbor (k-NN)**

- The k-NN prediction for  $\mathbf{x}$  is

$$f(\mathbf{x}) = \frac{1}{k} \sum_{\mathbf{x}_n \in nbh_k(\mathbf{x})} y_n$$

- where  $nbh_k(\mathbf{x})$  is the neighborhood of  $\mathbf{x}$  defined by the  $k$  closest points  $\mathbf{x}_n$ .
- **Curse of dimensionality:** Generalizing correctly becomes exponentially harder as the dimensionality grows.
- Gathering more inputs variables may be bad

12 **Support Vector Machine**

- Combination of the kernel trick plus a modified loss function (Hinge loss)
- Solution to the dual problem is sparse and non-zero entries will be our **support vectors**.
- **Kernelised feature vector** where  $\mu_k$  are centroids

$$\phi(\mathbf{x}) = [k(\mathbf{x}, \mu_1), \dots, k(\mathbf{x}, \mu_K)]$$

- In practice we'll take a subset of data points to be prototype  $\rightarrow$  **sparse vector machine**.
- Assume  $y_n \in \{-1, 1\}$
- SVM optimizes the following cost

$$\mathcal{L}(\mathbf{w}) = \min_{\mathbf{w}} \sum_{n=1}^N [1 - y_n \phi_n^T \mathbf{w}]_+ + \frac{\lambda}{2} \sum_{j=1}^M w_j^2$$

- Minimum doesn't change with a rescaling of  $\mathbf{w}$
- choose the hyperplane so that the distance from it to the nearest data point on each side is maximized
- **Duality:**

- Hard to minimize  $g(\mathbf{w})$  so we define

$$\mathcal{L}(\mathbf{w}) = \max_{\alpha} G(\mathbf{w}, \alpha)$$

- we use the property that  $C[v_n]_+ = \max(0, Cv_n) = \max_{\alpha_n \in [0, C]} \alpha_n v_n$

- We can rewrite the problem as

$$\min_{\alpha} \max_{\alpha \in [0, C]^N} \sum_{n=1}^N \alpha_n (1 - y_n \phi_n^T \mathbf{w}) + \frac{1}{2} \sum_{j=1}^M w_j^2$$

- This is differentiable, convex in  $\mathbf{w}$  and concave in  $\alpha$
- **Minimax theorem:**  
 $\min_{\mathbf{w}} \max_{\alpha} G(\mathbf{w}, \alpha) = \max_{\alpha} \min_{\mathbf{w}} G(\mathbf{w}, \alpha)$   
because  $G$  is convex in  $\mathbf{w}$  and concave in  $\alpha$ .
- Derivative w.r.t.  $\mathbf{w}$ :

$$\nabla_{\mathbf{w}} G(\mathbf{w}, \alpha) = - \sum_{n=1}^N \alpha_n y_n \mathbf{x}_n + \lambda \mathbf{w}$$

- Equating this to 0, we get:

$$\mathbf{w}(\alpha) = \frac{1}{\lambda} \sum_{n=1}^N \alpha_n y_n \mathbf{x}_n = \frac{1}{\lambda} \mathbf{X} \mathbf{Y} \alpha$$

$$\mathbf{Y} := \text{diag}(\mathbf{y})$$

- Plugging  $\mathbf{w}^*$  back in the dual problem  $\max_{\alpha \in [0, 1]^N} \alpha^T \mathbf{1} - \frac{1}{2\lambda} \alpha^T \mathbf{Y} \mathbf{X}^T \mathbf{X} \mathbf{Y} \alpha$
- This is a differentiable least-squares problem. Optimization is easy using Sequential Minimal Optimization. It is also naturally kernelized with  $\mathbf{K} = \mathbf{X}^T \mathbf{X}$
- The solution  $\alpha$  is sparse and is non-zero only for the training examples that are instrumental in determining the decision boundary.

13 **Kernel Ridge Regression**

- The following is true for ridge regression

$$\begin{aligned} \mathbf{w}^* &= (\mathbf{X}^T \mathbf{X} + \lambda \mathbf{I}_D)^{-1} \mathbf{X}^T \mathbf{y} \\ &= \mathbf{X}^T (\mathbf{X} \mathbf{X}^T + \lambda \mathbf{I}_N)^{-1} \mathbf{y} = \mathbf{X}^T \alpha^* \end{aligned} \tag{1}$$

- Complexity of computing  $\mathbf{w}$ : (1)  $O(D^2 N + D^3)$ , (2)  $O(DN^2 + N^3)$

- Thus we have

$$\mathbf{w}^* = \mathbf{X} \alpha^*, \quad \text{with } \mathbf{w}^* \in \mathbb{R}^D \text{ and } \alpha^* \in \mathbb{R}^N$$

- The representer theorem allows us to write an equivalent optimization problem in terms of  $\alpha$ .

$$\alpha = \underset{\alpha}{\text{argmax}} \left( -\frac{1}{2} \alpha^T (\mathbf{X}^T \mathbf{X} + \lambda \mathbf{I}_N) \alpha + \alpha^T \mathbf{y} \right)$$

- $\mathbf{K} = \mathbf{X} \mathbf{X}^T$  is called the **kernel matrix** or **Gram matrix**.
- If  $\mathbf{K}$  is positive definite, then it's called a **Mercer Kernel**.
- $\mathbf{K}_{i,j} = k(\mathbf{x}_i, \mathbf{x}_j)$
- If the kernel is Mercer, then there exists a function  $\phi(\mathbf{x})$  s.t.

$$k(\mathbf{x}, \mathbf{x}') = \phi(\mathbf{x})^T \phi(\mathbf{x}')$$

- **Kernel trick:**

- We can work directly with  $\mathbf{K}$  and never have to worry about  $\mathbf{X}$
- Replace  $(\mathbf{x}, \mathbf{x}')$  with  $k(\mathbf{x}, \mathbf{x}')$ .
- Kernel function can be interpreted as a measure of similarity
- The evaluation of a kernel is usually faster with  $k$  than with  $\phi$

- Kernelized ridge regression might be computationally more efficient in some cases.
- **Radial Basis function kernel (RBF)**

$$k(\mathbf{x}, \mathbf{x}') = \exp\left(-\frac{1}{2}(\mathbf{x} - \mathbf{x}')^T(\mathbf{x} - \mathbf{x}')\right)$$

- Properties of a kernel to ensure the existence of a corresponding  $\phi$ :
  - $\mathbf{K}$  should be symmetric:  
 $k(\mathbf{x}, \mathbf{x}') = k(\mathbf{x}', \mathbf{x})$
  - $\mathbf{K}$  should be positive semidefinite.
- Thus we get

$$\mathbf{y} = \mathbf{w}^T \mathbf{x} = \sum_{i=1}^K \alpha_i \mathbf{x}_i^T \mathbf{x} = \sum_{i=1}^K \alpha_i k(\mathbf{x}, \mathbf{x}_i)$$

14 **K-means**

- **Unsupervised learning:** Represent particular input patterns in a way that reflects the statistical structure of the overall collections of input patterns.
- **Cluster** are groups of points whose inter-point distances are small compared to the distances outside the cluster.

$$\min_{\mathbf{z}, \mu} \mathcal{L}(\mathbf{z}, \mu) = \sum_{k=1}^K \sum_{n=1}^N z_{nk} \|\mathbf{x}_n - \mu_k\|_2^2$$

- such that  $z_{nk} \in \{0, 1\}$  and  $\sum_{k=1}^K z_{nk} = 1$
- K-means algorithm (Coordinate Descent): Initialize  $\mu_k$ , then iterate

1. For all  $n$ , compute  $\mathbf{z}_n$  given  $\mu$   
$$z_{nk} = \begin{cases} 1 & \text{if } k = \underset{j}{\text{argmin}} \|\mathbf{x}_n - \mu_j\|_2^2 \\ 0 & \text{otherwise} \end{cases}$$
2. For all  $k$ , compute  $\mu_k$  given  $\mathbf{z}$   
$$\mu_k = \frac{\sum_{n=1}^N z_{nk} \mathbf{x}_n}{\sum_{n=1}^N z_{nk}}$$

- A good initialization procedure is to choose the prototypes to be equal to a random subset of  $K$  data points.
- Probabilistic model

$$p(\mathbf{z}, \mu) = \prod_{n=1}^N \prod_{k=1}^K [\mathcal{N}(\mathbf{x}_n | \mu_k, \mathbf{I})]^{z_{nk}}$$

- K-means as a Matrix Factorization  
$$\min_{\mathbf{z}, \mu} \mathcal{L}(\mathbf{z}, \mu) = \|\mathbf{X} - \mathbf{M} \mathbf{Z}^T\|_{\text{Frob}}^2$$
- Computation can be heavy, each example can belong to only on cluster and clusters have to be spherical.

15 **Gaussian Mixture Models**

- Clusters can be elliptical using a full covariance matrix instead of isotropic

covariance.

$$p(\mathbf{X} | \mu, \Sigma, \mathbf{z}) = \prod_{n=1}^N \prod_{k=1}^K [\mathcal{N}(\mathbf{x}_n | \mu_k, \Sigma_k)]^{z_{nk}}$$

- **Soft-clustering:** Points can belong to several cluster by defining  $z_n$  to be a random variable.

$$p(z_n = k) = \pi_k \text{ where } \pi_k > 0, \forall k, \sum_{k=1}^K \pi_k = 1$$

- Joint distribution of Gaussian mixture model

$$\begin{aligned} p(\mathbf{X}, \mathbf{z} | \mu, \Sigma, \pi) &= \prod_{n=1}^N p(\mathbf{x}_n | \mathbf{r}_n, \mu, \Sigma) p(\mathbf{z}_n | \pi) \\ &= \prod_{n=1}^N \prod_{k=1}^K [(\mathcal{N}(\mathbf{x}_n | \mu_k, \Sigma_k))^{z_{nk}}] \prod_{k=1}^K [\pi]^{z_{nk}} \end{aligned}$$

- $z_n$  are called *latent* unobserved variables
- Unknown parameters are given by  $\theta = \{\mu, \Sigma, \pi\}$
- We get the **marginal likelihood** by marginalizing  $z_n$  out from the likelihood

$$\begin{aligned} p(\mathbf{x}_n | \theta) &= \sum_{k=1}^K p(\mathbf{x}_n, z_n = k | \theta) \\ &= \sum_{k=1}^K p(z_n = k | \theta) p(\mathbf{x}_n | z_n = k, \theta) \\ &= \sum_{k=1}^K \pi_k \mathcal{N}(\mathbf{x}_n | \mu_k, \Sigma_k) \end{aligned}$$

- Without a latent variable model, number of parameters grow at rate  $O(N)$
- After marginalization, the growth is reduced to  $O(D^2 K)$
- To get maximum likelihood estimate of  $\theta$ , we maximize

$$\max_{\theta} \sum_{n=1}^N \log \sum_{k=1}^K \pi_k \mathcal{N}(\mathbf{x}_n | \mu_k, \Sigma_k)$$

16 **Expectation Maximization Algorithm**

- [ALGORITHM] Start with  $\theta^{(1)}$  and iterate

1. *Expectation step:* Compute a lower bound to the cost such that it is tight at the previous  $\theta^{(t)}$  with equality when,

$$q_{kn} = \frac{\pi_k \mathcal{N}(\mathbf{x}_n | \mu_k, \Sigma_k)}{\sum_{k=1}^K \pi_k \mathcal{N}(\mathbf{x}_n | \mu_k, \Sigma_k)}$$

2. *Maximization step:* Update  $\theta$   
$$\theta^{(t+1)} = \underset{\theta}{\text{argmax}} \mathcal{L}(\theta, \theta^{(t)})$$

$$\mu_k^{(t+1)} = \frac{\sum_{n=1}^N q_{kn}^{(i)} (r_{nk}) \mathbf{x}_n}{\sum_{n=1}^N q_{kn}^{(t)}}$$

$$\Sigma_k^{(t+1)} = \frac{\sum_{n=1}^N q_{kn}^{(t)} (\mathbf{x}_n - \mu_k^{(t+1)}) (\mathbf{x}_n - \mu_k^{(t+1)})^T}{\sum_{n=1}^N q_{kn}^{(t)}}$$

$$\pi_k^{(t+1)} = \frac{1}{N} \sum_{n=1}^N q_{kn}^{(t)}$$

- If covariance is diagonal  $\rightarrow$  K-means.

17 **Matrix factorization**

- We have  $D$  movies and  $N$  users
- $\mathbf{X}$  is a matrix  $D \times N$  with  $x_{dn}$  the rating of  $n$ 'th user for  $d$ 'th movie.
- We project data vectors  $\mathbf{x}_n$  to a smaller dimension  $\mathbf{z}_n \in \mathbb{R}^M$
- We have now 2 latent variables:
  - $\mathbf{Z}$  a  $N \times K$  matrix that gives features for the users
  - $\mathbf{W}$  a  $D \times K$  matrix that gives features for the movies

$$x_{dn} \approx \mathbf{w}_d^T \mathbf{z}_n$$

- We can add a regularizer and minimize the

following cost:

$$\begin{aligned} \mathcal{L}(\mathbf{W}, \mathbf{Z}) &= \frac{1}{2} \sum_{(d,n) \in \Omega} [x_{dn} - (\mathbf{W} \mathbf{Z}^T)_{dn}]^2 \\ &\quad + \frac{\lambda_w}{2} \sum_{d=1}^D \mathbf{w}_d^T \mathbf{w}_d + \frac{\lambda_z}{2} \sum_{n=1}^N \mathbf{z}_n^T \mathbf{z}_n \end{aligned}$$

- We can use coordinate descent algorithm, by first minimizing w.r.t.  $\mathbf{Z}$  given  $\mathbf{W}$  and then minimizing  $\mathbf{W}$  given  $\mathbf{Z}$ . This is called **Alternating least-squares (ALS)**:

$$\mathbf{Z}^T \leftarrow (\mathbf{W}^T \mathbf{W} + \lambda_z \mathbf{I}_K)^{-1} \mathbf{W}^T \mathbf{X}$$

$$\mathbf{W}^T \leftarrow (\mathbf{Z}^T \mathbf{Z} + \lambda_w \mathbf{I}_K)^{-1} \mathbf{Z}^T \mathbf{X}^T$$

- Complexity:  $O(DNK^2 + NK^3) \rightarrow O(DNK^2)$
- Probabilistic model

$$\begin{aligned} \prod_{(d,n) \in \Omega} \mathcal{N}(x_{dn} | \mathbf{w}_d^T \mathbf{z}_n, I) &\times \prod_{n=1}^N \mathcal{N}(\mathbf{z}_n | 0, \frac{1}{\lambda_z} I) \\ &\times \prod_{d=1}^D \mathcal{N}(\mathbf{w}_d | 0, \frac{1}{\lambda_w} I) \end{aligned}$$

- Since many ratings are missing we cannot normalize the data. A solution is to add offset terms:

$$\frac{1}{2} \sum_{(d,n) \in \Omega} (x_{dn} - \mathbf{w}_d^T \mathbf{z}_n - w_{0d} - z_{0n} - \mu)^2$$

18 **Singular Value Decomposition**

- Matrix factorization method

$$\mathbf{X} = \mathbf{U} \mathbf{S} \mathbf{V}^T$$

- $\mathbf{U}$  is a unitary  $D \times D$  matrix
- $\mathbf{V}$  is a unitary  $N \times N$  matrix
- $\mathbf{S}$  is a non-negative diagonal matrix of size  $D \times N$  which are called **singular values** appearing in a descending order.
- Columns of  $\mathbf{U}$  and  $\mathbf{V}$  are the left and right **singular vectors** respectively.
- Assuming  $D < N$  we have

$$\mathbf{X} = \sum_{d=1}^D s_d \mathbf{u}_d \mathbf{v}_d^T$$

This tells you about the spectrum of  $\mathbf{X}$  where higher singular vectors contain the *low-frequency information* and lower singular values contain the *high-frequency information*.

- Dimensionality Reduction

Take the matrix  $\mathbf{S}^{(K)}$  with the  $K$  first diagonal elements non zero. Then, rank- $K$  approx:

$$\mathbf{X} \approx \mathbf{X}_K = \mathbf{U} \mathbf{S}^{(K)} \mathbf{V}^T$$

18.1 **Principal Component Analysis**

PCA is a dimensionality reduction method and a method to decorrelate the data  $\mathbf{X} \approx \tilde{\mathbf{X}} = \mathbf{W} \mathbf{Z}^T$  such that columns of  $\mathbf{W}$  are orthogonal.

- If the data is zero mean

$$\Sigma = \frac{1}{N} \mathbf{X} \mathbf{X}^T \Rightarrow \mathbf{X} \mathbf{X}^T = \mathbf{U} \mathbf{S}^2 \mathbf{U}^T$$

$$\Rightarrow \mathbf{U}^T \mathbf{X} \mathbf{X}^T \mathbf{U} = \mathbf{U}^T \mathbf{U} \mathbf{S}^2 \mathbf{U}^T \mathbf{U} = \mathbf{S}^2$$

- Thus the columns of matrix  $\mathbf{U}$  are called the **principal components** and they decorrelate the covariance matrix.
- Using SVD, we can compute the matrices in the following way

$$\mathbf{W} = \mathbf{U} \mathbf{S}_D^{1/2}, \mathbf{Z}^T = \mathbf{S}^{1/2} \mathbf{V}^T$$

19 **Neural Net**

- Basic structure: One *input* layer of size  $D$ ,  $L$  *hidden* layers of size  $K$ , and one *output* layer. (*feedforward* network).

$$x_j^{(l)} = \phi \left( \sum_i w_{i,j}^{(l)} x_i^{(l-1)} + b_j^{(l)} \right).$$

- NN can represent the Riemann sum with only two layers  $\Rightarrow$  It's powerful!

- Cost function:

$$\begin{aligned} \mathcal{L} &= \frac{1}{N} \sum_{n=1}^N \left( y_n - f^{(L+1)} \circ \dots \circ f^{(1)}(\mathbf{x}_n^{(0)}) \right)^2 \\ \text{We can use SGD to minimize the cost function.} \\ \text{Compact form: } \mathbf{W}_{i,j}^{(l)} &= w_{i,j}^{(l)} \\ \mathbf{x}^{(l)} &= f^{(l)}(\mathbf{x}^{(l-1)}) = \\ &\phi \left( (\mathbf{W}^{(l)})^T \mathbf{x}^{(l-1)} + \mathbf{b}^{(l)} \right) \end{aligned}$$

19.1 **Backpropagation Algorithm**

- *Forward pass:* Compute

$$\mathbf{z}^{(l)} = (\mathbf{W}^{(l)})^T \mathbf{x}^{(l-1)} + \mathbf{b}^{(l)} \text{ with}$$

- $\mathbf{x}^{(0)} = \mathbf{x}_n$  and  $\mathbf{x}^{(l)} = \phi(\mathbf{z}^{(l)})$ .
- *Backward pass:* Set  $\delta^{(L+1)} = -2(y_n - \mathbf{x}^{(L+1)}) \phi'(z^{(L+1)})$  (if squared loss). Then compute

$$\begin{aligned} \delta_j^{(l)} &= \frac{\partial \mathcal{L}_n}{\partial z_j^{(l)}} = \sum_k \frac{\partial \mathcal{L}_n}{\partial z_k^{(l+1)}} \frac{\partial z_k^{(l+1)}}{\partial z_j^{(l)}} \\ &= \sum_k \delta_k^{(l+1)} \mathbf{W}_{j,k}^{(l+1)} \phi'(z_j^{(l)}) \end{aligned}$$

- *Final Computation:*

$$\frac{\partial \mathcal{L}_n}{\partial w_{i,j}^{(l)}} = \sum_k \frac{\partial \mathcal{L}_n}{\partial z_k^{(l)}} \frac{\partial z_k^{(l)}}{\partial w_{i,j}^{(l)}} = \frac{\partial \mathcal{L}_n}{\partial z_j^{(l)}} \frac{\partial z_j^{(l)}}{\partial w_{i,j}^{(l)}} = \delta_j^{(l)} \mathbf{x}_i^{(l-1)}$$

$$\frac{\partial \mathcal{L}_n}{\partial b_j^{(l)}} = \sum_k \frac{\partial \mathcal{L}_n}{\partial z_k^{(l)}} \frac{\partial z_k^{(l)}}{\partial b_j^{(l)}} = \frac{\partial \mathcal{L}_n}{\partial z_j^{(l)}} \frac{\partial z_j^{(l)}}{\partial b_j^{(l)}} = \delta_j^{(l)} \cdot 1 = \delta_j^{(l)}$$

19.2 **Activation Functions**

**Sigmoid**  $\phi(x) = \frac{1}{1+e^{-x}}$  Positive, bounded.

$\phi'(x) \simeq 0$  for large  $|x| \Rightarrow$  Learning slow.

**Tanh**  $\tanh(x) = \frac{e^x - e^{-x}}{e^x + e^{-x}} = \phi(2x) - 1/2$ .

Balanced, bounded. Learning slow too.

**ReLU**  $(x)_+ = \max(0, x)$  Positive, unbounded.

Derivate = 1 if  $x > 0$ , 0 if  $x < 0$

**Leaky ReLU**  $f(x) = \max \alpha x, x$  Remove 0 derivative.

**Maxout**

$$f(x) = \max \mathbf{x}^T \mathbf{w}_1 + b_1, \dots, \mathbf{x}^T \mathbf{w}_k + b_k$$

(Generalization of ReLU)

19.3 **Convolutional NN**

WHAT CAN I SAY???

19.4 **Reg, Data Augmentation and Dropout**

DO WE NEED SOMETHING IN THIS??

20 **Bayes Net**

- Graph example:  $p(x, y, z) = p(y|x)p(z|x)p(x)$   
: ( $y \leftarrow x \rightarrow z$ )

- **D-Separation**  $X$  and  $Y$  are D-separated by  $Z$  if every path from  $x \in X$  to  $y \in Y$  is blocked by  $Z$ .

- **Blocked Path** if the path contains a variable that

- is in  $Z$  and is **head-to-tail** or **tail-to-tail**
- the node is **head-to-head** and neither the node nor the descendant are in  $Z$ .

- **Markov Blanket** (which blocks node  $A$  from the rest of the net) contains:

- parents of  $A$
- children of  $A$
- parents of children of  $A$