

PCML Cheat Sheet

1 Math Prerequisites

- Bayes rule

$$p(A, B) = \underbrace{p(A|B)}_{\text{Lik.}} \underbrace{p(B)}_{\text{Prior}} = \underbrace{p(B|A)}_{\text{Post}} \underbrace{p(A)}_{\text{Marg. Lik.}}$$

- Bayes rule: $P(A|B) \propto P(B|A)P(A)$
- Gaussian distribution

$$\mathcal{N}(x|\mu, \sigma^2) = \frac{1}{\sqrt{2\pi\sigma^2}} \exp\left(-\frac{(x-\mu)^2}{2\sigma^2}\right)$$

$$\mathcal{N}(\mathbf{x}|\boldsymbol{\mu}, \boldsymbol{\Sigma}) = \frac{1}{\sqrt{2\pi|\boldsymbol{\Sigma}|}} \exp\left(-\frac{1}{2}(\mathbf{x}-\boldsymbol{\mu})^T \boldsymbol{\Sigma}^{-1}(\mathbf{x}-\boldsymbol{\mu})\right)$$

- Production of independent variables:
 $\text{Var}(XY) = \mathbb{E}(X^2) \mathbb{E}(Y^2) - [\mathbb{E}(X)]^2 [\mathbb{E}(Y)]^2$
- Covariance matrix of a data vector \mathbf{x}

$$\boldsymbol{\Sigma} = \frac{1}{N} \sum_{n=1}^N (\mathbf{x}_n - \mathbb{E}(\mathbf{x}))(\mathbf{x}_n - \mathbb{E}(\mathbf{x}))^T$$

1.1 Convexity

- The Hessian of a convex function is psd and for a strictly-convex function it's pd.

$$\mathbf{H}_{i,j} = d^2 f / dx_i dx_j$$

1.2 Linear Algebra

- Condition number** If \mathbf{A} is normal ($\mathbf{A}^T \mathbf{A} = \mathbf{A} \mathbf{A}^T$) then

$$k(\mathbf{A}) = \left| \frac{\lambda_{\max}(\mathbf{A})}{\lambda_{\min}(\mathbf{A})} \right|$$

2 Cost functions

Mean square error (MSE):

$$MSE(\mathbf{w}) = \sum_{n=1}^N (y_n - f(\mathbf{x}_n))^2$$

Mean Absolute Error (MAE):

$$MAE = \sum_{n=1}^N |y_n - f(\mathbf{x}_n)|$$

Huber loss

$$Huber = \begin{cases} \frac{1}{2} z^2 & , |z| \leq \delta \\ \delta |z| - \frac{1}{2} \delta^2 & , |z| > \delta \end{cases}$$

Hinge loss

$$Hinge = [1 - y_n f(\mathbf{x}_n)]_+ = \max(0, 1 - y_n f(\mathbf{x}_n))$$

Logistic loss

$$Logistic = \log(1 - \exp(y_n f(\mathbf{x}_n)))$$

3 Optimization

3.1 Grid search

- Compute the cost over a grid of M points to find the minimum. Exponential Complexity. Hard to find a good range of values

3.2 Gradient Descent

- GD uses only first-order information and takes steps in the opposite direction of the gradient
- Given cost function $\mathcal{L}(\mathbf{w})$ we want to find $\mathbf{w} = \arg \min_{\mathbf{w}} \mathcal{L}(\mathbf{w})$

3.3 Batch Gradient Descent

- Take steps in the opposite direction of the gradient

$$\mathbf{w}^{(t+1)} \leftarrow \mathbf{w}^{(t)} - \gamma \nabla \mathcal{L}(\mathbf{w}^{(t)})$$

- with $\gamma > 0$ the learning rate.
- With γ too big, method might diverge. With γ too small, convergence is slow.

3.4 Gradients for MSE

- We define the error vector \mathbf{e} :

$$\mathbf{e} := \mathbf{y} - \mathbf{X}\mathbf{w}$$

- and MSE as follows:

$$\mathcal{L}(\mathbf{w}) = \frac{1}{2N} \sum_{n=1}^N (y_n - \tilde{\mathbf{x}}_n^T \mathbf{w})^2 = \frac{1}{2N} \mathbf{e}^T \mathbf{e}$$

- then the gradient is given by

$$\nabla \mathcal{L}(\mathbf{w}) = -\frac{1}{N} \mathbf{X}^T \mathbf{e}$$

- Optimality conditions:

- necessary*: gradient equal zero:
 $\frac{d\mathcal{L}(\mathbf{w}^*)}{d\mathbf{w}} = 0$
- sufficient*: Hessian matrix is positive

- definite: $\mathbf{H}(\mathbf{w}^*) = \frac{d^2 \mathcal{L}(\mathbf{w}^*)}{d\mathbf{w} d\mathbf{w}^T}$
- Very sensitive to illconditioning \Rightarrow always normalize features.
- Complexity*: $O(NDI)$ with I the number of iterations

3.5 Stochastic Gradient Descent

In ML, most cost functions are formulated as a sum over the training examples:

$$\mathcal{L}(\mathbf{w}) = \frac{1}{N} \sum_{n=1}^N \mathcal{L}_n(\mathbf{w})$$

\Rightarrow SGD algo is given by update rule:

$$\mathbf{w}^{(t+1)} \leftarrow \mathbf{w}^{(t)} - \gamma \nabla \mathcal{L}_n(\mathbf{w}^{(t)})$$

Idea: Cheap but unbiased estimate of grad.

$$\mathbb{E}[\nabla \mathcal{L}_n(\mathbf{w})] = \nabla \mathcal{L}(\mathbf{w})$$

3.6 Mini-batch SGD

Update direction ($B \subseteq [N]$):

$$\mathbf{g}^{(t)} := \frac{1}{|B|} \sum_{n \in B} \nabla \mathcal{L}_n(\mathbf{w}^{(t)})$$

Update rule: $\mathbf{w}^{(t+1)} \leftarrow \mathbf{w}^{(t)} - \gamma \mathbf{g}^{(t)}$

3.7 Subgradients (Non-Smooth optim)

A vector $\mathbf{g} \in \mathbb{R}^D$ s.t.

$$\mathcal{L}(\mathbf{u}) \geq \mathcal{L}(\mathbf{w}) + \mathbf{g}^T (\mathbf{u} - \mathbf{w}) \quad \forall \mathbf{u}$$

is the subgradient to \mathcal{L} at \mathbf{w} . If \mathcal{L} is differentiable at \mathbf{w} , we have $\mathbf{g} = \nabla \mathcal{L}(\mathbf{w})$

4 Least Squares

- Use the first optimality conditions:

$$\nabla \mathcal{L}(\mathbf{w}^*) = 0 \Rightarrow \mathbf{X}^T \mathbf{e} = \mathbf{X}^T (\mathbf{y} - \mathbf{X}\mathbf{w}) = 0$$

- When $\mathbf{X}^T \mathbf{X}$ is invertible, we have the closed-form expression

$$\mathbf{w}^* = (\mathbf{X}^T \mathbf{X})^{-1} \mathbf{X}^T \mathbf{y}$$

- thus we can predict values for a new \mathbf{x}_m

$$\mathbf{y}_m := \mathbf{x}_m^T \mathbf{w}^* = \mathbf{x}_m^T (\mathbf{X}^T \mathbf{X})^{-1} \mathbf{X}^T \mathbf{y}$$

- The **Gram matrix** $\mathbf{X}^T \mathbf{X}$ is pd and is also invertible iff \mathbf{X} has full column rank.
- Complexity*: $O(ND^2 + D^3) \equiv O(ND^2)$
- \mathbf{X} can be rank deficient when $D > N$ or when the columns $\tilde{\mathbf{x}}_d$ are nearly collinear. \Rightarrow matrix is ill-conditioned.

5 Maximum Likelihood

- Let define our mistakes $\epsilon_n \sim \mathcal{N}(0, \sigma^2)$.

$$\rightarrow y_n = \mathbf{x}_n^T \mathbf{w} + \epsilon_n$$

- Another way of expressing this:

$$p(\mathbf{y}|\mathbf{X}, \mathbf{w}) = \prod_{n=1}^N p(y_n|\mathbf{x}_n, \mathbf{w}) \\ = \prod_{n=1}^N \mathcal{N}(y_n|\mathbf{x}_n^T \mathbf{w}, \sigma^2)$$

which defines the likelihood of observing \mathbf{y} given \mathbf{X} and \mathbf{w}

- Define cost with log-likelihood

$$\mathcal{L}_{lik}(\mathbf{w}) = \log p(\mathbf{y}|\mathbf{X}, \mathbf{w})$$

$$= -\frac{1}{2\sigma^2} \sum_{n=1}^N (y_n - \mathbf{x}_n^T \mathbf{w})^2 + cnst$$

- Maximum likelihood estimator (MLE) gives another way to design cost functions

$$\underset{\mathbf{w}}{\text{argmin}} \mathcal{L}_{MSE}(\mathbf{w}) = \underset{\mathbf{w}}{\text{argmax}} \mathcal{L}_{lik}(\mathbf{w})$$

- MLE can also be interpreted as finding the model under which the observed data is most likely to have been generated from.
- With Laplace distribution

$$p(y_n|\mathbf{x}_n, \mathbf{w}) = \frac{1}{2b} e^{-\frac{1}{b}|y_n - \mathbf{x}_n^T \mathbf{w}|}$$

$$\sum_n \log p(y_n|\mathbf{x}_n, \mathbf{w}) = \sum_n |y_n - \mathbf{x}_n^T \mathbf{w}| + cnst$$

6 Ridge Regression

- Basis functions:

$$y_n = w_0 + \sum_{j=1}^M w_j \phi_j(\mathbf{x}_n) = \tilde{\boldsymbol{\phi}}(\mathbf{x}_n)^T \mathbf{w}$$

- This model is linear in \mathbf{w} but nonlinear in \mathbf{x} . Dimension is now M, not N.
- Polynomial basis

$$\phi(x_n) = [1, x_n, x_n^2, \dots, x_n^M]$$

- The least square solution becomes

$$\mathbf{w}_{lse}^* = (\tilde{\boldsymbol{\Phi}}^T \tilde{\boldsymbol{\Phi}})^{-1} \tilde{\boldsymbol{\Phi}}^T \mathbf{y}$$

- Complex models overfit easily. Thus we can penalize them with a **regularization term**

$$\min_{\mathbf{w}} \left(\mathcal{L}(\mathbf{w}) + \frac{\lambda}{2N} \sum_{j=1}^M w_j^2 \right)$$

$$\mathbf{w}^* = \underset{\mathbf{w}}{\text{argmin}} \left(\frac{1}{2} (\mathbf{y} - \mathbf{X}\mathbf{w})^T (\mathbf{y} - \mathbf{X}\mathbf{w}) + \frac{\lambda}{2} \mathbf{w}^T \mathbf{w} \right)$$

- Note that w_0 is not penalized.
- By differentiating and setting to zero we get

$$\mathbf{w}_{ridge} = (\tilde{\boldsymbol{\Phi}}^T \tilde{\boldsymbol{\Phi}} + \Lambda)^{-1} \tilde{\boldsymbol{\Phi}}^T \mathbf{y}$$

$$\Lambda = \begin{bmatrix} 0 & \lambda I_M \\ 0 & \lambda I_M \end{bmatrix}$$

- Ridge regression improves the condition number of the Gram matrix since the eigenvalues of $(\tilde{\boldsymbol{\Phi}}^T \tilde{\boldsymbol{\Phi}} + \lambda I_m)$ are at least λ :
SVD $\tilde{\boldsymbol{\Phi}}^T \tilde{\boldsymbol{\Phi}} = U S U^T$ with $S = \text{diag}(\geq 0)$, then $\tilde{\boldsymbol{\Phi}}^T \tilde{\boldsymbol{\Phi}} + \Lambda = U(S + \Lambda)U^T$ has eigenvalues $\geq \lambda$.

- Maximum-a-posteriori (MAP) estimator**:

- Maximizes the product of the likelihood and the prior.

$$\mathbf{w}_{MAP} = \underset{\mathbf{w}}{\text{argmax}} p(\mathbf{y}|\mathbf{X}, \mathbf{w}) p(\mathbf{w}|\boldsymbol{\Sigma})$$

- Assume $w_0 = 0$

$$\mathbf{w}_{ridge} = \underset{\mathbf{w}}{\text{argmax}} \left(\log \left[\prod_{n=1}^N \mathcal{N}(y_n|\mathbf{x}_n^T \mathbf{w}, \Lambda) \times \mathcal{N}(\mathbf{w}|\mathbf{0}, \mathbf{I}) \right] \right)$$

- Lasso regularizer** forces some w_i to be strictly 0 and therefore forces sparsity in the model.

$$\min_{\mathbf{w}} \frac{1}{2N} \sum_{n=1}^N (y_n - \tilde{\boldsymbol{\phi}}(\mathbf{x}_n)^T \mathbf{w})^2,$$

$$\text{such that } \sum_{i=1}^M |w_i| \leq \tau$$

7 Bias-Variance decomposition

- The expected test error can be expressed as the sum of two terms

$$\text{expected loss} = (\text{bias})^2 + \text{variance} + \text{noise}$$

- Model bias and estimation bias are important
- RR increases estimation bias and reduces var
- Model more complex increases test error

- Small $\lambda \rightarrow$ low bias but large variance
- Large $\lambda \rightarrow$ large bias but low variance

$$err = \sigma^2 + \mathbb{E}[f_{lse} - \mathbb{E}[f_{lse}]]^2 + [f_{true} - \mathbb{E}[f_{lse}]]^2$$

8 Logistic Regression

- Classification** relates input variables \mathbf{x} to discrete output variable y
- Binary classifier**: we use $y = 0$ for \mathbf{C}_1 and $y = 1$ for \mathbf{C}_2 .
- Can use least-squares to predict \hat{y}_*

$$\hat{y} = \begin{cases} \mathbf{C}_1 & \hat{y}_* < 0.5 \\ \mathbf{C}_2 & \hat{y}_* \geq 0.5 \end{cases}$$

- Logistic function**

$$\sigma(x) = \frac{\exp(x)}{1 + \exp(x)}$$

$$p(y_n = \mathbf{C}_1|\mathbf{x}_n) = \sigma(\mathbf{x}_n^T \mathbf{w})$$

$$p(y_n = \mathbf{C}_2|\mathbf{x}_n) = 1 - \sigma(\mathbf{x}_n^T \mathbf{w})$$

- The probabilistic model:

$$p(\mathbf{y}|\mathbf{X}, \mathbf{w}) = \prod_{n=1}^N \sigma(\mathbf{x}_n^T \mathbf{w})^{y_n} (1 - \sigma(\mathbf{x}_n^T \mathbf{w}))^{1-y_n}$$

- The log-likelihood:

$$\mathcal{L}_{MLE}(\mathbf{w}) = \sum_{n=1}^N (y_n \mathbf{x}_n^T \mathbf{w} - \log(1 + \exp(\mathbf{x}_n^T \mathbf{w})))$$

- We can use the fact that

$$\frac{d}{dx} \log(1 + \exp(x)) = \sigma(x)$$

- Gradient of the log-likelihood

$$\mathbf{g} = \frac{d\mathcal{L}}{d\mathbf{w}} = \sum_{n=1}^N (\mathbf{x}_n y_n - \mathbf{x}_n \sigma(\mathbf{x}_n^T \mathbf{w}))$$

$$= -\mathbf{X}^T [\sigma(\mathbf{X}\mathbf{w}) - \mathbf{y}]$$

- The negative of the log-likelihood $-\mathcal{L}_{MLE}(\mathbf{w})$ is convex
- Hessian** of the log-likelihood

- We know that $\frac{d\sigma(t)}{dt} = \sigma(t)(1 - \sigma(t))$

- Hessian is the derivative of the gradient

$$\mathbf{H}(\mathbf{w}) = -\frac{d\mathbf{g}(\mathbf{w})}{d\mathbf{w}^T} = \sum_{n=1}^N \frac{d}{d\mathbf{w}^T} \sigma(\mathbf{x}_n^T \mathbf{w}) \mathbf{x}_n$$

$$= \sum_{n=1}^N \mathbf{x}_n \sigma(\mathbf{x}_n^T \mathbf{w}) (1 - \sigma(\mathbf{x}_n^T \mathbf{w})) \mathbf{x}_n^T$$

$$= \tilde{\mathbf{X}}^T \mathbf{S} \tilde{\mathbf{X}}$$

where \mathbf{S} is a $N \times N$ diagonal matrix with diagonals

$$S_{nn} = \sigma(\mathbf{x}_n^T \mathbf{w}) (1 - \sigma(\mathbf{x}_n^T \mathbf{w}))$$

- The negative of the log-likelihood is not strictly convex.

Newton's Method

- Uses second-order information and takes steps in the direction that minimizes a quadratic approximation

$$\mathcal{L}(\mathbf{w}) = \mathcal{L}(\mathbf{w}^{(k)}) + \nabla \mathcal{L}_k^T (\mathbf{w} - \mathbf{w}^{(k)})$$

$$+ \frac{1}{2} (\mathbf{w} - \mathbf{w}^{(k)})^T \mathbf{H}_k (\mathbf{w} - \mathbf{w}^{(k)})$$

so

$$\nabla \mathcal{L}(\mathbf{w}) = \nabla \mathcal{L}(\mathbf{w}^{(k)}) + \mathbf{H}_k (\mathbf{w} - \mathbf{w}^{(k)})$$

equating to 0 gives the minimum:

$$\mathbf{w}^{k+1} = \mathbf{w}^{(k)} - \gamma_k \mathbf{H}_k^{-1} \nabla \mathcal{L}_k$$

- γ_k added because it's all approximative.

- Complexity*: $O((ND^2 + D^3)I)$

Penalized Logistic Regression

$$\min_{\mathbf{w}} \left(-\sum_{n=1}^N \log p(y_n|\mathbf{x}_n^T \mathbf{w}) + \lambda \sum_{d=1}^D w_d^2 \right)$$

9 Generalized Linear Model

Exponential family distribution

$$p(\mathbf{y}|\boldsymbol{\eta}) = h(y) \exp(\boldsymbol{\eta}^T \boldsymbol{\phi}(\mathbf{y}) - A(\boldsymbol{\eta}))$$

$$A(\boldsymbol{\eta}) = \ln \left[\int_y h(y) e^{\boldsymbol{\eta}^T \boldsymbol{\phi}(y)} dy \right] < \infty$$

is convex by Holder.

- Bernoulli distribution

$$p(y|\mu) = \mu^y (1 - \mu)^{1-y}$$

$$= \exp(y \log(\frac{\mu}{1-\mu}) + \log(1-\mu))$$

- there is a relationship between η and μ through the **link function**

$$\eta = \log\left(\frac{\mu}{1-\mu}\right) \leftrightarrow \mu = \frac{e^\eta}{1+e^\eta}$$

- Note that μ is the mean parameter of y

- Relationship between the mean $\boldsymbol{\mu}$ and $\boldsymbol{\eta}$ is defined using a link function g

$$\boldsymbol{\eta} = \mathbf{g}(\boldsymbol{\mu}) \Leftrightarrow \boldsymbol{\mu} = \mathbf{g}^{-1}(\boldsymbol{\eta})$$

- First and second derivatives of $A(\boldsymbol{\eta})$ are related to the mean and the variance

$$\frac{dA(\boldsymbol{\eta})}{d\boldsymbol{\eta}} = \mathbb{E}[\boldsymbol{\phi}(\boldsymbol{\eta})], \quad \frac{d^2 A(\boldsymbol{\eta})}{d\boldsymbol{\eta}^2} = \text{Var}[\boldsymbol{\phi}(\boldsymbol{\eta})]$$

- $A(\boldsymbol{\eta})$ is convex

- The generalized maximum likelihood cost to minimize is

$$\min_{\mathbf{w}} \mathcal{L}(\mathbf{w}) = -\sum_{n=1}^N \log(p(y_n|\mathbf{x}_n^T \mathbf{w}))$$

where $p(y_n|\mathbf{x}_n^T \mathbf{w})$ is an exponential family distribution

- We obtain the solution

$$\frac{d\mathcal{L}}{d\mathbf{w}} = \mathbf{X}^T [\mathbf{g}^{-1}(\mathbf{X}\mathbf{w}) - \boldsymbol{\phi}(\mathbf{y})]$$

since $\frac{dA(\boldsymbol{\eta})}{d\boldsymbol{\eta}} = \mathbf{g}^{-1}(\boldsymbol{\eta})$.

10 k-Nearest Neighbor (k-NN)

- The k-NN prediction for \mathbf{x} is

$$f(\mathbf{x}) = \frac{1}{k} \sum_{\mathbf{x}_n \in nbh_k(\mathbf{x})} y_n$$

where $nbh_k(\mathbf{x})$ is the neighborhood of \mathbf{x} defined by the k closest points \mathbf{x}_n . For classification take majority vote instead of mean.

- Curse of dimensionality**: Let $\mathcal{X} = [0, 1]^d$, $\mathcal{Y} = \{0, 1\}$, $S \sim \mathcal{D}^N$, $\eta(\mathbf{x}) = p(y = 1 | \mathbf{x})$. If \mathcal{D} is known, best classifier is Bayes $f_*(\mathbf{x}) = 1[\eta(\mathbf{x}) > \frac{1}{2}]$. If

$$|\eta(\mathbf{x}) - \eta(\mathbf{x}')| \leq c \|\mathbf{x} - \mathbf{x}'\|$$

- **Duality:**
 - Hard to minimize $g(\mathbf{w})$ so we define $\mathcal{L}(\mathbf{w}) = \max_{\alpha} G(\mathbf{w}, \alpha)$
 - we use the property that $C[v_n]_+ = \max(0, C v_n) = \max_{\alpha_n \in [0, C]} \alpha_n v_n$
 - We can rewrite the problem as
$$\min_{\mathbf{w}} \max_{\alpha \in [0, C]^N} \sum_{n=1}^N \alpha_n (1 - y_n \phi_n^T \mathbf{w}) + \frac{1}{2} \sum_{j=1}^M w_j^2$$
 - This is differentiable, convex in \mathbf{w} and concave in α
 - **Minimax theorem:** $\min_{\mathbf{w}} \max_{\alpha} G(\mathbf{w}, \alpha) = \max_{\mathbf{w}} \min_{\alpha} G(\mathbf{w}, \alpha)$ because G is convex in \mathbf{w} and concave in α .
 - Derivative w.r.t. \mathbf{w} :
$$\nabla_{\mathbf{w}} G(\mathbf{w}, \alpha) = - \sum_{n=1}^N \alpha_n y_n \mathbf{x}_n + \lambda \mathbf{w}$$
 - Equating this to 0, we get:
$$\mathbf{w}(\alpha) = \frac{1}{\lambda} \sum_{n=1}^N \alpha_n y_n \mathbf{x}_n = \frac{1}{\lambda} \mathbf{X} \mathbf{Y} \alpha$$

$$\mathbf{Y} := \text{diag}(\mathbf{y})$$
 - Plugging \mathbf{w}^* back in the dual problem
$$\max_{\alpha \in [0, 1]^N} \alpha^T \mathbf{1} - \frac{1}{2\lambda} \alpha^T \mathbf{Y}^T \mathbf{X}^T \mathbf{X} \mathbf{Y} \alpha$$
 - This is a differentiable least-squares problem. Optimization is easy using Sequential Minimal Optimization. It is also naturally kernelized with $\mathbf{K} = \mathbf{X}^T \mathbf{X}$
 - The solution α is sparse and is non-zero only for the training examples that are instrumental in determining the decision boundary.
 - $\alpha_n = 0$ if x_n is far from the boundary in the right side. $\alpha_n \in (0, 1)$ if x_n is exactly on the margin and $\alpha_n = 1$ if x_n is in the margin or on the wrong side.
- 12 Kernel Ridge Regression**
 - The following is true for ridge regression
$$\mathbf{w}^* = (\mathbf{X}^T \mathbf{X} + \lambda \mathbf{I}_D)^{-1} \mathbf{X}^T \mathbf{y}$$

$$= \mathbf{X}^T (\mathbf{X} \mathbf{X}^T + \lambda \mathbf{I}_N)^{-1} \mathbf{y} = \mathbf{X}^T \alpha^*$$
 - since $(PQ + I_N)^{-1} P = P(QP + I_M)^{-1}$.
 - Complexity of computing \mathbf{w} : (1) $O(D^2 N + D^3)$, (2) $O(DN^2 + N^3)$
 - Thus we have $\mathbf{w}^* = \mathbf{X} \alpha^*$, with $\mathbf{w}^* \in \mathbb{R}^D$ and $\alpha^* \in \mathbb{R}^N$
 - The representer theorem allows us to write an equivalent optimization problem in terms of α .
$$\alpha = \arg \max_{\alpha} \left(-\frac{1}{2} \alpha^T (\mathbf{X}^T \mathbf{X} + \lambda \mathbf{I}_N) \alpha + \alpha^T \mathbf{y} \right)$$
 - $\mathbf{K} = \mathbf{X} \mathbf{X}^T$ is called the **kernel matrix** or **Gram matrix**.
 - If \mathbf{K} is positive definite, then it's called a **Mercer Kernel**.
 - $\mathbf{K}_{i,j} = k(\mathbf{x}_i, \mathbf{x}_j)$
 - If the kernel is Mercer, then there exists a function $\phi(\mathbf{x})$ s.t.
$$k(\mathbf{x}, \mathbf{x}') = \phi(\mathbf{x})^T \phi(\mathbf{x}')$$
 - **Kernel trick:**
 - compute dot-product in \mathbb{R}^m while remaining in \mathbb{R}^n
 - Replace $\langle \mathbf{x}, \mathbf{x}' \rangle$ with $k(\mathbf{x}, \mathbf{x}')$.
 - **Common Kernel**
 - Polynomial Kernel: $(\gamma \langle \mathbf{x}_i, \mathbf{x}_j \rangle + r)^d$
 - Radial Basis function kernel (RBF)
$$k(\mathbf{x}, \mathbf{x}') = \exp\left(-\frac{1}{2}(\mathbf{x} - \mathbf{x}')^T (\mathbf{x} - \mathbf{x}')\right)$$
 - Sigmoid Kernel: $\tanh(\langle \mathbf{x}_i, \mathbf{x}_j \rangle + r)$

- **Properties of kernels** to ensure the existence of a corresponding ϕ :
 - symmetric: $k(\mathbf{x}, \mathbf{x}') = k(\mathbf{x}', \mathbf{x})$
 - positive semi-definite.
- Thus we get
$$\mathbf{y} = \mathbf{w}^T \mathbf{x} = \sum_{i=1}^K \alpha_i \mathbf{x}_i^T \mathbf{x} = \sum_{i=1}^K \alpha_i k(\mathbf{x}, \mathbf{x}_i)$$
- 13 K-means**
 - **Unsupervised learning:** Represent particular input patterns in a way that reflects the statistical structure of the overall collections of input patterns.
 - **Cluster** are groups of points whose inter-point distances are small compared to the distances outside the cluster.
$$\min_{\mathbf{z}, \mu} \mathcal{L}(\mathbf{z}, \mu) = \sum_{k=1}^K \sum_{n=1}^N z_{nk} \|\mathbf{x}_n - \mu_k\|_2^2$$
 such that $z_{nk} \in \{0, 1\}$ and $\sum_{n=1}^N z_{nk} = 1$
 - K-means algorithm (Coordinate Descent): Initialize μ_k , then iterate
 - For all n , compute μ_k given μ
$$z_{nk} = \begin{cases} 1 & \text{if } k = \arg \min_j \|\mathbf{x}_n - \mu_j\|_2^2 \\ 0 & \text{otherwise} \end{cases}$$
 - For all k , compute μ_k given \mathbf{z}
$$\mu_k = \frac{\sum_{n=1}^N z_{nk} \mathbf{x}_n}{\sum_{n=1}^N z_{nk}}$$
 - A good initialization procedure is to choose the prototypes to be equal to a random subset of K data points.
 - Probabilistic model
$$p(\mathbf{z}, \mu) = \prod_{n=1}^N \prod_{k=1}^K [\mathcal{N}(\mathbf{x}_n | \mu_k, \mathbf{I})]^{z_{nk}}$$
 - K-means as a Matrix Factorization
$$\min_{\mathbf{z}, \mu} \mathcal{L}(\mathbf{z}, \mu) = \|\mathbf{X} - \mathbf{M} \mathbf{Z}^T\|_{\text{Frob}}^2$$
 - Computation can be heavy, each example can belong to only on cluster and clusters have to be spherical.
- 14 Gaussian Mixture Models**
 - Clusters can be elliptical using a full covariance matrix instead of isotropic covariance.
$$p(\mathbf{X} | \mu, \Sigma, \mathbf{z}) = \prod_{n=1}^N \prod_{k=1}^K [\mathcal{N}(\mathbf{x}_n | \mu_k, \Sigma_k)]^{z_{nk}}$$
 - **Soft-clustering:** Points can belong to several cluster by defining z_n to be a random variable.
 - Joint distribution of Gaussian mixture model
$$p(\mathbf{X}, \mathbf{z} | \mu, \Sigma, \pi) = \prod_{n=1}^N p(\mathbf{x}_n | \mathbf{r}_n, \mu, \Sigma) p(\mathbf{z}_n | \pi)$$

$$= \prod_{n=1}^N \prod_{k=1}^K [(\mathcal{N}(\mathbf{x}_n | \mu_k, \Sigma_k))^{z_{nk}}] \prod_{k=1}^K [\pi]^{z_{nk}}$$
 - z_n are called *latent* unobserved variables
 - Unknown parameters are given by $\theta = \{\mu, \Sigma, \pi\}$
 - We get the **marginal likelihood** by marginalizing z_n out from the likelihood
$$p(\mathbf{x}_n | \theta) = \sum_{k=1}^K p(\mathbf{x}_n, z_n = k | \theta)$$

$$= \sum_{k=1}^K p(z_n = k | \theta) p(\mathbf{x}_n | z_n = k, \theta)$$

$$= \sum_{k=1}^K \pi_k \mathcal{N}(\mathbf{x}_n | \mu_k, \Sigma_k)$$

- Without a latent variable model, number of parameters grow at rate $O(N)$
- After marginalization, the growth is reduced to $O(D^2 K)$
- To get maximum likelihood estimate of θ , we maximize
$$\max_{\theta} \sum_{n=1}^N \log \sum_{k=1}^K \pi_k \mathcal{N}(\mathbf{x}_n | \mu_k, \Sigma_k)$$
- 15 Expectation Maximization Algorithm**
 - *[ALGORITHM]* Start with $\theta^{(1)}$ and iterate
 - Expectation step:* Compute a lower bound to the cost such that it is tight at the previous $\theta^{(t)}$ with equality when,
$$q_{kn} = \frac{\pi_k \mathcal{N}(\mathbf{x}_n | \mu_k, \Sigma_k)}{\sum_{k=1}^K \pi_k \mathcal{N}(\mathbf{x}_n | \mu_k, \Sigma_k)}$$
 this is the posterior of $z_n = k$ given \mathbf{x}_n, θ .
 - Maximization step:* Update θ
$$\theta^{(t+1)} = \arg \max_{\theta} \mathcal{L}(\theta, \theta^{(t)})$$

$$\mu_k^{(t+1)} = \frac{\sum_{n=1}^N \gamma^{(t)}(r_{nk}) \mathbf{x}_n}{\sum_{n=1}^N q_{kn}^{(t)}}$$

$$\Sigma_k^{(t+1)} = \frac{\sum_{n=1}^N q_{kn}^{(t)} (\mathbf{x}_n - \mu_k^{(t+1)}) (\mathbf{x}_n - \mu_k^{(t+1)})^T}{\sum_{n=1}^N q_{kn}^{(t)}}$$

$$\pi_k^{(t+1)} = \frac{1}{N} \sum_{n=1}^N q_{kn}^{(t)}$$
 - If covariance is diagonal \rightarrow K-means.
- 16 Matrix factorization**
 - We have D movies and N users
 - \mathbf{X} is a matrix $D \times N$ with x_{dn} the rating of n 'th user for d 'th movie.
 - We project data vectors \mathbf{x}_n to a smaller dimension $\mathbf{z}_n \in \mathbb{R}^M$
 - We have now 2 latent variables:
 - \mathbf{Z} a $N \times K$ matrix that gives features for the users
 - \mathbf{W} a $D \times K$ matrix that gives features for the movies
 - Non-convex and non-identifiable ($W = \beta W, Z = \frac{1}{\beta} Z$).
 - We can add a regularizer and minimize the following cost:
$$\mathcal{L}(\mathbf{W}, \mathbf{Z}) = \frac{1}{2} \sum_{(d,n) \in \Omega} [x_{dn} - (\mathbf{W} \mathbf{Z}^T)_{dn}]^2 + \frac{\lambda_w}{2} \sum_{d=1}^D \mathbf{w}_d^T \mathbf{w}_d + \frac{\lambda_z}{2} \sum_{n=1}^N \mathbf{z}_n^T \mathbf{z}_n$$
 - We can use coordinate descent algorithm, by first minimizing w.r.t. \mathbf{Z} given \mathbf{W} and then minimizing \mathbf{W} given \mathbf{Z} . This is called **Alternating least-squares (ALS)**:
$$\mathbf{Z}^T \leftarrow (\mathbf{W}^T \mathbf{W} + \lambda_z \mathbf{I}_K)^{-1} \mathbf{W}^T \mathbf{X}$$

$$\mathbf{W}^T \leftarrow (\mathbf{Z}^T \mathbf{Z} + \lambda_w \mathbf{I}_K)^{-1} \mathbf{Z}^T \mathbf{X}^T$$
 - *Complexity:* $O(DNK^2 + NK^3) \rightarrow O(DNK^2)$
 - Probabilistic model
$$\prod_{(d,n) \in \Omega} \mathcal{N}(x_{dn} | \mathbf{w}_d^T \mathbf{z}_n, I) \times \prod_{n=1}^N \mathcal{N}(\mathbf{z}_n | 0, \frac{1}{\lambda_z} I)$$

$$\times \prod_{d=1}^D \mathcal{N}(\mathbf{w}_d | 0, \frac{1}{\lambda_w} I)$$
 - Since many ratings are missing we cannot normalize the data. A solution is to add offset terms:
$$\frac{1}{2} \sum_{(d,n) \in \Omega} (x_{dn} - \mathbf{w}_d^T \mathbf{z}_n - w_{0d} - z_{0n} - \mu)^2$$

- 17 Singular Value Decomposition**
 - Matrix factorization method
$$\mathbf{X} = \mathbf{U} \mathbf{S} \mathbf{V}^T$$
 - \mathbf{U} is a unitary $D \times D$ matrix
 - \mathbf{V} is a unitary $N \times N$ matrix
 - \mathbf{S} is a non-negative diagonal matrix of size $D \times N$ which are called **singular values** appearing in a descending order.
 - Columns of \mathbf{U} and \mathbf{V} are the left and right **singular vectors** respectively.
 - Assuming $D < N$ we have
$$\mathbf{X} = \sum_{d=1}^D s_d \mathbf{u}_d \mathbf{v}_d^T$$
 - This tells you about the spectrum of \mathbf{X} where higher singular vectors contain the *low-frequency information* and lower singular values contain the *high-frequency information*.
 - Dimensionality Reduction
 - Take the matrix $\mathbf{S}^{(K)}$ with the K first diagonal elements non zero. Then, rank- K approx:
$$\mathbf{X} \approx \mathbf{X}_K = \mathbf{U} \mathbf{S}^{(K)} \mathbf{V}^T$$
 - 17.1 Principal Component Analysis**
 - PCA is a dimensionality reduction method and a method to decorrelate the data $\mathbf{X} \approx \tilde{\mathbf{X}} = \mathbf{W} \mathbf{Z}^T$ such that columns of \mathbf{W} are orthogonal.
 - If the data is zero mean
$$\Sigma = \frac{1}{N} \mathbf{X} \mathbf{X}^T \Rightarrow \mathbf{X} \mathbf{X}^T = \mathbf{U} \mathbf{S}^2 \mathbf{U}^T$$

$$\Rightarrow \mathbf{U}^T \mathbf{X} \mathbf{X}^T \mathbf{U} = \mathbf{U}^T \mathbf{U} \mathbf{S}^2 \mathbf{U}^T \mathbf{U} = \mathbf{S}^2$$
 - Thus the columns of matrix \mathbf{U} are called the **principal components** and they decorrelate the covariance matrix.
 - Using SVD, we can compute the matrices in the following way
$$\mathbf{W} = \mathbf{U} \mathbf{S}_D^{1/2}, \mathbf{Z}^T = \mathbf{S}^{1/2} \mathbf{V}^T$$
 - 18 Neural Net**
 - Basic structure: One *input* layer of size D , L *hidden* layers of size K , and one *output* layer. (*feedforward* network).
$$x_j^{(l)} = \phi \left(\sum_i w_{i,j}^{(l-1)} x_i^{(l-1)} + b_j^{(l)} \right).$$
 - Lemma: For $f: \mathbb{R}^D \rightarrow \mathbb{R}$ with $\int |w| |f(w)| dw \leq C$ with
$$\tilde{f}(w) = \int f(x) e^{-i w^T x} dx$$
 is the Fourier transform. Then $\forall n > 0$, $\exists f_n(x) = \sum_{i=1}^n c_i \sigma(x^i w_i + b_i) + c_0$ so that
$$\int_{|x| < r} (f - f_n)^2 \leq \frac{(2Cr)^2}{n},$$
 with σ sigmoid-like.
 - Cost function:
$$\mathcal{L} = \frac{1}{N} \sum_{n=1}^N \left(y_n - f^{(L+1)} \circ \dots \circ f^{(1)}(\mathbf{x}_n^{(0)}) \right)^2$$
 - We can use SGD to minimize the cost function.
 - Compact form: $\mathbf{W}_{i,j}^{(l)} = w_{i,j}^{(l)}$
$$\mathbf{x}^{(l)} = f^{(l)}(\mathbf{x}^{(l-1)}) = \phi \left(\left(\mathbf{W}^{(l)} \right)^T \mathbf{x}^{(l-1)} + \mathbf{b}^{(l)} \right)$$
 - 18.1 Backpropagation Algorithm**
 - *Forward pass:* Compute
$$\mathbf{z}^{(l)} = \left(\mathbf{W}^{(l)} \right)^T \mathbf{x}^{(l-1)} + \mathbf{b}^{(l)}$$
 with
$$\mathbf{x}^{(0)} = \mathbf{x}_n$$
 and $\mathbf{x}^{(l)} = \phi(\mathbf{z}^{(l)})$.
 - *Backward pass:* Set
$$\delta^{(L+1)} = -2(y_n - \mathbf{x}^{(L+1)}) \phi'(z^{(L+1)})$$
 (if squared loss). Then compute

- $$\delta_j^{(l)} = \frac{\partial \mathcal{L}_n}{\partial z_j^{(l)}} = \sum_k \frac{\partial \mathcal{L}_n}{\partial z_k^{(l+1)}} \frac{\partial z_k^{(l+1)}}{\partial z_j^{(l)}} = \sum_k \delta_k^{(l+1)} \mathbf{W}_{j,k}^{(l+1)} \phi'(z_j^{(l)})$$
- *Final Computation:*
$$\frac{\partial \mathcal{L}_n}{\partial w_{i,j}^{(l)}} = \sum_k \frac{\partial \mathcal{L}_n}{\partial z_k^{(l)}} \frac{\partial z_k^{(l)}}{\partial w_{i,j}^{(l)}} = \frac{\partial \mathcal{L}_n}{\partial z_j^{(l)}} \frac{\partial z_j^{(l)}}{\partial w_{i,j}^{(l)}} = \delta_j^{(l)} \mathbf{x}_i^{(l-1)}$$

$$\frac{\partial \mathcal{L}_n}{\partial b_j^{(l)}} = \sum_k \frac{\partial \mathcal{L}_n}{\partial z_k^{(l)}} \frac{\partial z_k^{(l)}}{\partial b_j^{(l)}} = \frac{\partial \mathcal{L}_n}{\partial z_j^{(l)}} \frac{\partial z_j^{(l)}}{\partial b_j^{(l)}} = \delta_j^{(l)} \cdot 1 = \delta_j^{(l)}$$
- 18.2 Activation Functions**
 - Sigmoid** $\phi(x) = \frac{1}{1+e^{-x}}$ Positive, bounded. $\phi'(x) \simeq 0$ for large $|x| \Rightarrow$ Learning slow.
 - Tanh** $\tanh(x) = \frac{e^x - e^{-x}}{e^x + e^{-x}} = 2\phi(2x) - 1$. Balanced, bounded. Learning slow too.
 - ReLU** $(x)_+ = \max\{0, x\}$ Positive, unbounded. Derivate = 1 if $x > 0$, 0 if $x < 0$
 - Leaky ReLU** $f(x) = \max\{\alpha x, x\}$ Remove 0 derivative.
- Maxout**
$$f(x) = \max\{\mathbf{x}^T \mathbf{w}_1 + b_1, \dots, \mathbf{x}^T \mathbf{w}_k + b_k\}$$
 (Generalization of ReLU)
- 18.3 Convolutional NN**
 - Sparse connections and *weights sharing*: reduce complexity. (e.g. pixels in pictures only depend on neighbours)
- 18.4 Reg, Data Augmentation and Dropout**
 - Regularization term: $\frac{1}{2} \sum_{l=1}^L \mu^{(l)} \|\mathbf{W}^{(l)}\|_F^2$
 - Data Augm.: e.g. shift or rotation of pics
 - Dropout: Avoid overfit. Drop nodes randomly. (Then average multiple drop-NN)
- 19 Bayes Net**
 - Graph example: $p(x, y, z) = p(y|x)p(z|x)p(x)$: ($y \leftarrow x \rightarrow z$)
 - Lemma: X and Y are independent given Z iff X and Y are D -separated by Z .
 - **D-Separation** X and Y are D -separated by Z if every path from $x \in X$ to $y \in Y$ is blocked by Z .
 - **Blocked Path** contains a variable that
 - is in Z and is **head-to-tail** ($x \rightarrow z \rightarrow y$) or **tail-to-tail** ($x \leftarrow z \rightarrow y$).
 - the node is **head-to-head** ($x \rightarrow a \leftarrow y$) and neither the node nor the descendant are in Z .
 - **Markov Blanket** (which blocks node A from the rest of the net) contains:
 - parents of A
 - children of A
 - parents of children of A
 - Lemma: Any x is independent of $y \notin MB(X)$ given $MB(X)$.
 - $f(x_1, \dots, x_6) = f_1(x_1, x_2, x_3) f_2(x_1, x_4, x_6) f_3(x_4) f_4(x_4, x_5)$, then $f(x_1) = [\sum_{x_2, x_3} f_1(x_1, x_2, x_3)] [\sum_{x_4} f_3(x_4) (\sum_{x_6} f_2(x_1, x_4, x_6))]$